

Numerical Simulation and Scientific Computing II

Lecture 2: Distributed Parallel Computing II



Josef Weinbub, Paul Manstetten,
Heinz Pettermann, Asur Vijaya Kumar, Pavan
Kumar, Jesús Carrete Montana, Francesco
Zonta, Kevin Sturm

Institute for Microelectronics
TU Wien



nssc@iue.tuwien.ac.at

Quiz

Q1: How is it ensured that a specific message is received by a specific process?

→ tag and dest parameter

Q2: What is the first and last routine to be called in a MPI program?

→ *MPI_Init (...)* and *MPI_Finalize()*

Q3: Is “MPI_Init” executed by one, several or all MPI processes?

→ all

Q4: Name typical reduction operations?

→ sum, min, max, etc.

Q5: How can a point-to-point communication be made non-blocking? What potential advantage is there?

→ *MPI_I...*, potentially allows to overlap communication with computation

Outline


- Distributed Parallel Computing II
 - **Non-Blocking Communication**
 - Collective Communication
 - Derived Datatypes
 - Virtual Topologies
 - Hybrid MPI+OpenMP
- Quiz

Non-Blocking Communications


- **Separate communication into three phases:**
 - **Initiate nonblocking communication**
 - returns **I**mmediately
 - routine name starting with **MPI_I**...
 - **Do some work (perhaps involving other communications?)**
 - If cluster hardware doesn't have dedicated MPI communication logic, don't expect efficient overlapping of communication with computation: Communication requires significant computational resources!
 - **Wait for non-blocking communication to complete**

Non-Blocking Examples

- Non-blocking send




- Non-blocking receive




Non-Blocking Send

- **Initiate non-blocking send**
 - Ring example: Initiate non-blocking send to the right neighbor
→
- **Do some work**
 - Ring example: Receiving the message from left neighbor
→
- **Message transfer can be completed**
- **Wait for non-blocking send to complete** /



Non-Blocking Receive

- **Initiate non-blocking receive**
 - Ring example: Initiate non-blocking receive from left neighbor
→
- **Do some work**
 - Ring example: Sending the message to the right neighbor
→
- **Message transfer can be completed**
- **Wait for non-blocking receive to complete** /



Request Handles

- Request handles
 - Are used for non-blocking communication
 - Must be stored in local variables: MPI_Request
 - is generated by a non-blocking communication routine
 - is used (and freed) in the MPI_WAIT routine

Non-Blocking Synchronous Send

- buf must not be modified between Issend and Wait
- “Issend + Wait directly after” is equivalent to blocking call (Ssend)
- status is not used in Issend, but in Wait (with send: nothing returned)

```
MPI_Issend( buf, count, datatype, dest, tag, comm, [OUT] &request_handle);
```



```
MPI_Wait( [INOUT] &request_handle, &status);
```

Nonblocking Receive

- **buf must not be used between Irecv and Wait**

```
MPI_Irecv ( buf, count, datatype, source, tag, comm, [OUT] &request_handle);
```



```
MPI_Wait( [INOUT] &request_handle, &status);
```

Blocking and Non-Blocking

- **Send and receive can be blocking or non-blocking.**
- **A blocking send can be used with a non-blocking receive, and vice-versa.**
- **Non-blocking sends can use any mode**
 - **standard – MPI_ISEND**
 - **synchronous – MPI_ISSEND**
 - **buffered – MPI_IBSEND**
 - **ready – MPI_IRSEND**
- **Synchronous mode affects completion, i.e. MPI_Wait / MPI_Test, not initiation, i.e., MPI_I...**
- **The non-blocking operation immediately followed by a matching wait is equivalent to the blocking operation.**

Completion


```
MPI_Wait( &request_handle, &status);  
MPI_Test( &request_handle, &flag, &status);
```

- **You need to**
 - **WAIT or**
 - **Loop with TEST until request is completed, i.e., flag == 1**

Multiple Non-Blocking Communications

You have several request handles:

- **Wait or test for completion of one message**
 - `MPI_Waitany` / `MPI_Testany`
- **Wait or test for completion of all messages**
 - `MPI_Waitall` / `MPI_Testall` *)
- **Wait or test for completion of as many messages as possible (i.e., at least one)**
 - `MPI_Waitsome` / `MPI_Testsome` *)



*) Each status contains an additional error field.

This field is only used if `MPI_ERR_IN_STATUS` is returned (also valid for send operations).

Send-Receive in One Routine

- **MPI_Sendrecv & MPI_Sendrecv_replace**
 - Combines the triple “MPI_Irecv + Send + Wait” into one routine

Performance Options

Which is the fastest neighbor communication?

- **MPI_Irecv + MPI_Send**
- **MPI_Irecv + MPI_Isend**
- **MPI_Isend + MPI_Recv**
- **MPI_Isend + MPI_Irecv**
- **MPI_Sendrecv**
- **MPI_Neighbor_alltoall**

No answer by the MPI standard, because:

MPI targets portable and efficient message-passing programming but efficiency of MPI application-programming is not portable!

Example Non-Blocking Communication

- **Integration with MPI non-blocking communications**
 - **Example source:**
http://www.bu.edu/tech/support/research/training-consulting/online-tutorials/mpi/example1-2/example1_3/
 - **Background on numerical integration:**
<http://www.bu.edu/tech/support/research/training-consulting/online-tutorials/mpi/example1-2/>
- **Until a matching receive has signaled that it is ready to receive, a blocking send will continue to wait.**

```
void other_work(int myid) {  
    printf("more work on process %dn", myid);  
}  
float integral(float ai, float h, int n){  
    int j;  
    float aij, integ;  
    integ = 0.0;           /* initialize */  
    for (j=0;j<j++) {      /* sum integrals */  
        aij = ai + (j+0.5)*h; /* mid-point */  
        integ += cos(aij)*h;  
    }  
    return integ;  
}
```

**Support
Functions**

Example Non-Blocking Communication

Common Part

```
int n, p, myid, tag, master, proc, ierr;
float h, integral_sum, a, b, ai, pi, my_int;
MPI_Comm comm;
MPI_Request request;
MPI_Status status;
comm = MPI_COMM_WORLD;
ierr = MPI_Init(&argc, &argv);      /* starts MPI */
MPI_Comm_rank(comm, &myid);         /* get current process id */
MPI_Comm_size(comm, &p);            /* get number of processes */
master = 0;
pi = acos(-1.0); /* = 3.14159... */
a = 0.;          /* lower limit of integration */
b = pi*1./2.;    /* upper limit of integration */
n = 500;         /* number of increment within each process */
tag = 123;       /* set the tag to identify this particular job */
h = (b-a)/n/p;   /* length of increment */
ai = a + myid*n*h; /* lower limit of integration for partition myid */
my_int = integral(ai, h, n); /* 0<=myid<=p-1 */
printf("Process %d has the partial result of %fn", myid, my_int);
```

Example Non-Blocking Communication

Master Part

```
if(myid == master) {  
    integral_sum = my_int;  
    for (proc=1;proc<p;proc++) {  
        MPI_Recv(  
            &my_int, 1, MPI_FLOAT, /* triplet of buffer, size, data type */  
            MPI_ANY_SOURCE, /* message source */  
            MPI_ANY_TAG, /* message tag */  
            comm, &status); /* status identifies source, tag */  
        integral_sum += my_int;  
    }  
    printf("The Integral =%fn",integral_sum); /* sum of my_int */  
}
```

Example Non-Blocking Communication

Worker Part

```
else {  
    MPI_Isend(    /* non-blocking send */  
        &my_int, 1, MPI_FLOAT,    /* triplet of buffer, size, data type */  
        master,  
        tag,  
        comm, &request);    /* send my_int to master */  
    other_work(myid);  
    MPI_Wait(&request, &status);    /* block until Isend is done */  
}  
MPI_Finalize();    /* let MPI finish up ... */  
}
```

Outline

- Distributed Parallel Computing II
 - Non-Blocking Communication
 - **Collective Communication**
 - Derived Datatypes
 - Virtual Topologies
 - Hybrid MPI+OpenMP
- Quiz

Collective Communication

- **Communications involving a group of processes.**
- **Called by all processes in a communicator.**
- **Examples:**
 - **Barrier synchronization.**
 - **Broadcast, scatter, gather.**
 - **Global sum, global maximum, etc.**
 - **Neighbor communication in a virtual grid**

Characteristics of Collective Communication

- **Collective action over a communicator.**
- **All processes of the communicator must communicate, i.e., must call the collective routine.**
- **Synchronization may or may not occur, therefore all processes must be able to start the collective routine.**
- **On a given communicator, the n-th collective call must match on all processes of the communicator.**
- **In MPI-1.0 – MPI-2.2, all collective operations are blocking. Non-blocking versions since MPI-3.0.**
 - **Not covered in this lecture**
- **No tags.**
- **Receive buffers must have exactly the same size as send buffers.**


Barrier Synchronization

MPI_Barrier is normally never needed:

- **all synchronization is done automatically by the data communication:
a process cannot continue before it has the data that it needs.**
- **if used for debugging:
please guarantee, that it is removed in production.**
- **for profiling: to separate time measurements**

Broadcast

```
int MPI_Bcast(void *buf, int count, MPI_Datatype datatype,  
             int root, MPI_Comm comm)
```

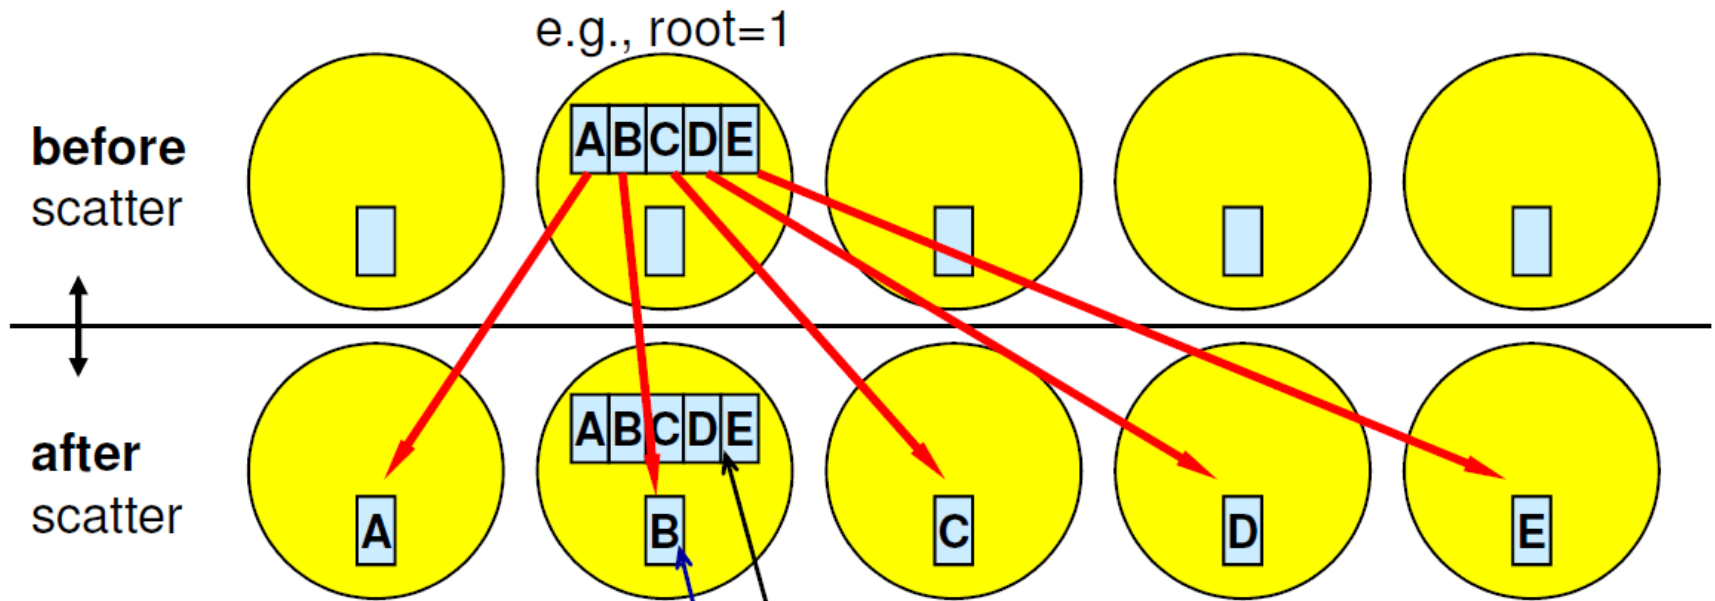


- E.g. root=1: all processes must use same "root"

Example

```
std::vector<char> buf(3);  
if(myrank == 1) { buf[0]='r'; buf[1]='e'; buf[2]='d'; }  
MPI_Bcast(buf.data(), buf.size(), MPI_CHAR, 1, MPI_COMM_WORLD);
```


Scatter



```
int MPI_Scatter(void *sendbuf, int sendcount, MPI_Datatype sendtype,  
               void *recvbuf, int recvcount, MPI_Datatype recvtype,  
               int root, MPI_Comm comm)
```

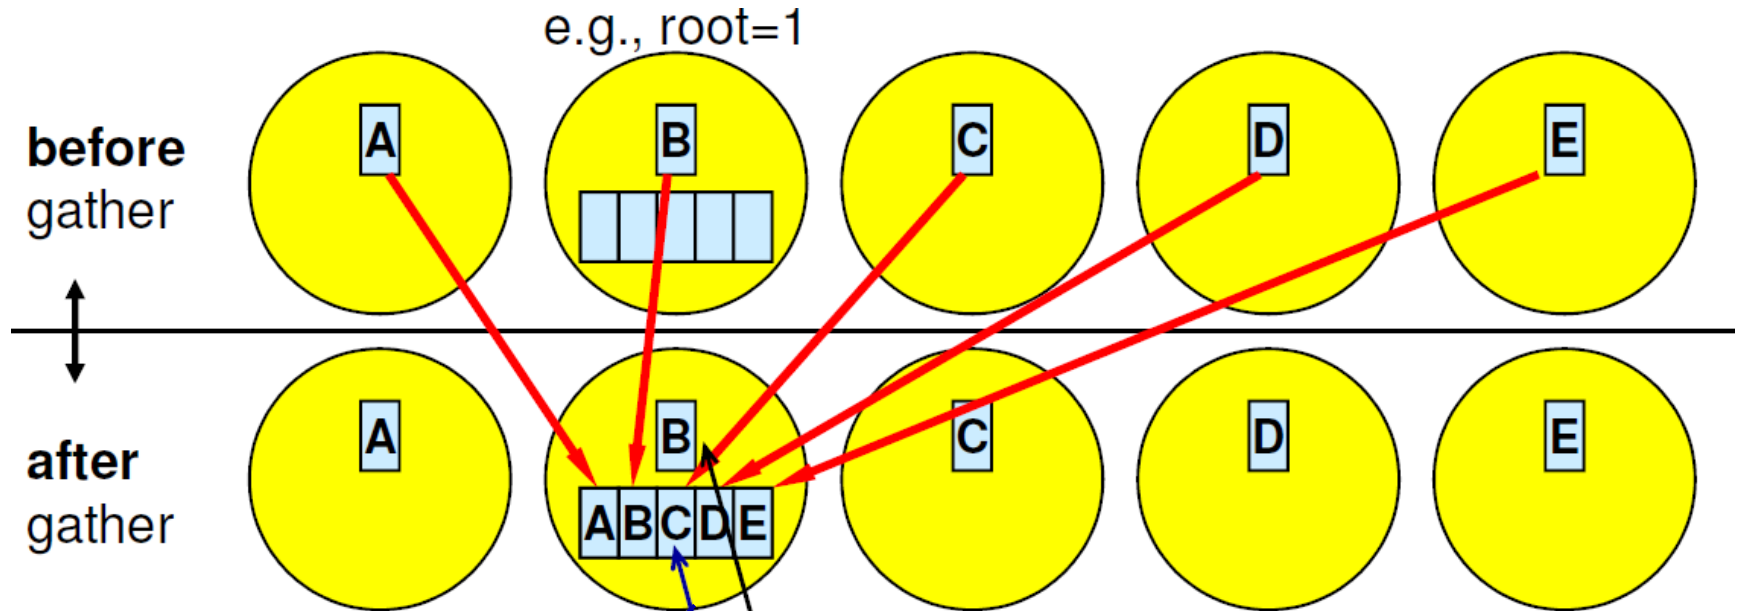
Example

```
std::vector<char> sbuf(5); char rbuf;  
if(myrank == 1) sbuf[0]='A'; sbuf[1]='B'; // etc.  
MPI_Scatter(sbuf.data(), 1, MPI_CHAR,  
            &rbuf, 1, MPI_CHAR,  
            1, MPI_COMM_WORLD)
```

If data cannot be equally distributed to all processes, use:
MPI_Scatterv: allows for variable send/recvcounts.

Sum of transmitted data size must match:
 $\text{sendcount} \times \text{sendtype} = \text{recvcount} \times \text{recvtype}$
In principle: allowed to mix data types!
But typically: counts and types the same.

Gather



```
int MPI_Gather(void *sendbuf, int sendcount, MPI_Datatype sendtype,  
              void *recvbuf, int recvcount, MPI_Datatype recvtype,  
              int root, MPI_Comm comm)
```

Example

```
if(myrank==0) char sbuf='A'; // etc... 'B', 'C', 'D', 'E'  
std::vector<char> rbuf(5);  
MPI_Gather(&sbuf, 1, MPI_CHAR,  
          rbuf.data(), 1, MPI_CHAR,  
          1, MPI_COMM_WORLD)
```

See notes for Scatter, same here.

Global Reduction Operations

- To perform a global reduce operation across
- $d_0 \circ d_1 \circ d_2 \circ d_3 \circ \dots \circ d_{s-2} \circ d_{s-1}$
 - d_i = data in process rank i
 - single variable, or
 - vector
 - \circ = associative operation
 - Example:
 - global sum or product
 - global maximum or minimum
 - global user-defined operation

Predefined Reduction Operation Handles

Predefined operation handle	Function
MPI_MAX	Maximum
MPI_MIN	Minimum
MPI_SUM	Sum
MPI_PROD	Product
MPI_LAND	Logical AND
MPI_BAND	Bitwise AND
MPI_LOR	Logical OR
MPI_BOR	Bitwise OR
MPI_LXOR	Logical exclusive OR
MPI_BXOR	Bitwise exclusive OR
MPI_MAXLOC	Maximum and location of the maximum
MPI_MINLOC	Minimum and location of the minimum

Reduce


```
int MPI_Reduce(void *inbuf, void *result, int count, MPI_Datatype datatype,  
              MPI_Op op, int root, MPI_Comm comm)
```

```
int inbuf = 5; // 2, 7, 4 for other ranks
```

```
int result;
```

```
MPI_Reduce(&inbuf, &result, 1, MPI_INT,  
          MPI_SUM, 0, MPI_COMM_WORLD)
```

MPI_Reduce




```
std::vector<int> inbuf {5,1}; // {2,3} ... for other ranks
```

```
std::vector<int> result (2);
```

```
MPI_Reduce(inbuf.data(), result.data(), result.size(),  
          MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD)
```

MPI_Reduce



Example Collective Communication

- Integration with MPI collective communications
 - Example source:
http://www.bu.edu/tech/support/research/training-consulting/online-tutorials/mpi/example1-2/example1_4/
 - Adaption of previous example based on non-blocking communication

```
float integral(float ai, float h, int n){  
    int j;  
    float aij, integ;  
    integ = 0.0;           /* initialize */  
    for (j=0;j<j++;) {     /* sum integrals */  
        aij = ai + (j+0.5)*h; /* mid-point */  
        integ += cos(aij)*h;  
    }  
    return integ;  
}
```

**Support
Function**

Example Collective Communication

First Part

```
int n, p, myid, tag, proc, ierr, i;
float h, integral_sum, a, b, ai, pi, my_int, buf[50];
int master = 0; /* processor performing total sum */
MPI_Comm comm;
comm = MPI_COMM_WORLD;
ierr = MPI_Init(&argc,&argv); /* starts MPI */
MPI_Comm_rank(comm, &myid); /* get current process id */
MPI_Comm_size(comm, &p); /* get number of processes */
pi = acos(-1.0); /* = 3.14159... */
a = 0.; /* lower limit of integration */
b = pi*1./2.; /* upper limit of integration */
n = 500; /* number of increment within each process */
tag = 123; /* set the tag to identify this particular job */
h = (b-a)/n/p; /* length of increment */
ai = a + myid*n*h; /* lower limit of integration for partition myid */
my_int = integral(ai, h, n); /* 0<=myid<=p-1 */
printf("Process %d has the partial sum of %fn", myid, my_int);
```

buf-size \geq p
(see next slide)

Example Collective Communication

```
MPI_Gather( /* collects my_int from all processes to master */  
    &my_int, 1, MPI_FLOAT, /* send buffer, size, data type */  
    &buf[0], 1, MPI_FLOAT, /* receive buffer, size, data type */  
    master, comm);
```

Second Part

```
if(myid == master) {  
    integral_sum = 0.0;  
    for (i=0; i<p; i++) {  
        integral_sum += buf[i];  
    }  
    printf("The Integral =%fn",integral_sum);  
}
```

buf must fit all ranks!

```
MPI_Finalize();
```

Is there another collective operation which could be used here?

Yes, MPI_Reduce with MPI_SUM

Outline

- Distributed Parallel Computing II
 - Non-Blocking Communication
 - Collective Communication
 - **Derived Datatypes**
 - Virtual Topologies
 - Hybrid MPI+OpenMP
- Quiz

Derived Datatypes

- **Description of the memory layout of the buffer**
 - for sending
 - for receiving
- **Basic types**
- **Derived types**
 - vectors
 - structs
 - others

Data Layout and the Describing Datatype Handle

```
struct buff_layout
```

```
{ int    i_val[3];  
  double d_val[5];  
} buffer;
```

Compiler

```
array_of_types[0]=MPI_INT;  
array_of_blocklengths[0]=3;  
array_of_displacements[0]=  
array_of_types[1]=MPI_DOUBLE;  
array_of_blocklengths[1]=5;  
array_of_displacements[1]=
```


```
MPI_Type_create_struct(2, array_of_blocklengths,  
                      array_of_displacements, array_of_types,  
                      &buff_datatype);
```

```
MPI_Type_commit(&buff_datatype);
```

```
MPI_Send(&buffer, 1, buff_datatype, ...)
```

&buffer = the start
address of the data

the datatype handle
describes the data layout



Compute Displacement

```
array_of_types[0]=MPI_INT;  
array_of_blocklengths[0]=3;  
array_of_displacements[0]=  
array_of_types[1]=MPI_DOUBLE;  
array_of_blocklengths[1]=5;  
array_of_displacements[1]=  
MPI_Type_create_struct(2, array_of_blocklengths,  
array_of_displacements, array_of_types,  
                        &buff_datatype);  
MPI_Type_commit(&buff_datatype);
```

In principle:

`array_of_displacements[i] = address(block_i) – address(block_0);`

Use:

`int MPI_Get_address(void* location, MPI_Aint *address)`

Continuing the example:

```
int address_base, address_0, address_1;  
MPI_Get_address(&buffer, &address_base);  
MPI_Get_address(&buffer.i_val[0], &address_0);  
MPI_Get_address(&buffer.d_val[0], &address_1);  
array_of_displacements[0] = address_0 – address_base;  
array_of_displacements[1] = address_1 – address_base;
```


```
struct buff_layout  
{ int      i_val[3];  
  double d_val[5];  
} buffer;
```

Derived Datatype: Another Example

```
typedef struct {  
    char    a;  
    int     b;  
    double  c;} mystruct;  
mystruct mydata;  
int baseaddr, addr0, addr1, addr2;  
MPI_Get_address ( &mydata, &baseaddr);  
MPI_Get_address ( &mydata.a, &addr0);  
MPI_Get_address ( &mydata.b, &addr1);  
MPI_Get_address ( &mydata.c, &addr2);  
displ[0] = addr0 - baseaddr;  
displ[1] = addr1 - baseaddr;  
displ[2] = addr2 - baseaddr;  
dtype[0] = MPI_CHAR;  
length[0] = 1;  
dtype[1] = MPI_INT;  
length[1] = 1;  
dtype[2] = MPI_DOUBLE;  
length[2] = 1;  
MPI_Type_struct ( 3, length, displ, dtype, &newtype );  
MPI_Type_commit ( &newtype );
```

Contiguous Data

- The simplest derived datatype
- Consists of a number of contiguous items of the same datatype



```
int MPI_Type_contiguous(int count, MPI_Datatype oldtype,  
                        MPI_Datatype *newtype)
```

Example

```
int myvec[4];  
MPI_Type_contiguous ( 4, MPI_INT, &mybrandnewdatatype);  
MPI_Type_commit ( &mybrandnewdatatype );  
MPI_Send ( myvec, 1, mybrandnewdatatype, ... );
```

Committing and Freeing a Datatype

- Before a datatype handle is used in message passing communication, it needs to be committed with **MPI_TYPE_COMMIT**.
- This needs to be done only once (by each MPI process).
- If type no longer required, one may call **MPI_TYPE_FREE()** to free a datatype and its internal resources.


```
int MPI_Type_commit(MPI_Datatype *datatype);  
int MPI_Type_free  (MPI_Datatype *datatype);
```

Outline


- Distributed Parallel Computing II
 - Non-Blocking Communication
 - Collective Communication
 - Derived Datatypes
 - **Virtual Topologies**
 - Hybrid MPI+OpenMP
- Quiz

Virtual Topologies

- Convenient process naming.
- Naming scheme to fit the communication pattern.
- Simplifies writing of code.
- Can allow MPI to optimize communications.
- Normal topology:



- Now consider that you work with this:




How to Use a Virtual Topology

- **Creating a topology produces a new communicator.**
- **MPI provides mapping functions:**
 - **to compute process ranks, based on the topology naming scheme,**
 - **and vice versa.**

Example – A Two-Dimensional Cylinder

- Ranks and Cartesian process coordinates



Topology Types

- **Cartesian Topologies**

- each process is connected to its neighbor in a virtual grid,
- boundaries can be cyclic, or not,
- processes are identified by Cartesian coordinates,
- of course, communication between any two processes is still allowed.

- **Graph Topologies**

- general graphs,
- two interfaces:
 - `MPI_GRAPH_CREATE` (since MPI-1)
 - `MPI_DIST_GRAPH_CREATE_ADJACENT` &
 - `MPI_DIST_GRAPH_CREATE` (new scalable interface since MPI-2.2)
- not covered here.

Creating a Cartesian Virtual Topology

```
int MPI_Cart_create(MPI_Comm comm_old, int ndims,  
                   int *dims, int *periods, int reorder,  
                   MPI_Comm *comm_cart)
```

- **Comm_old** = **MPI_COMM_WORLD**

- **ndims** = **2**

- **Dims** = **(4, 3)**


- **Periods** = **(1, 0)**

- **Reorder** = **0 or 1**

→ Reorder: Task MPI backend to optimally assign MPI ranks to specific Cartesian coordinates


considering communication ramifications!

→ But then ranks in comm_cart differ from comm_old!




Cartesian Mapping Functions

- Mapping ranks to process grid coordinates



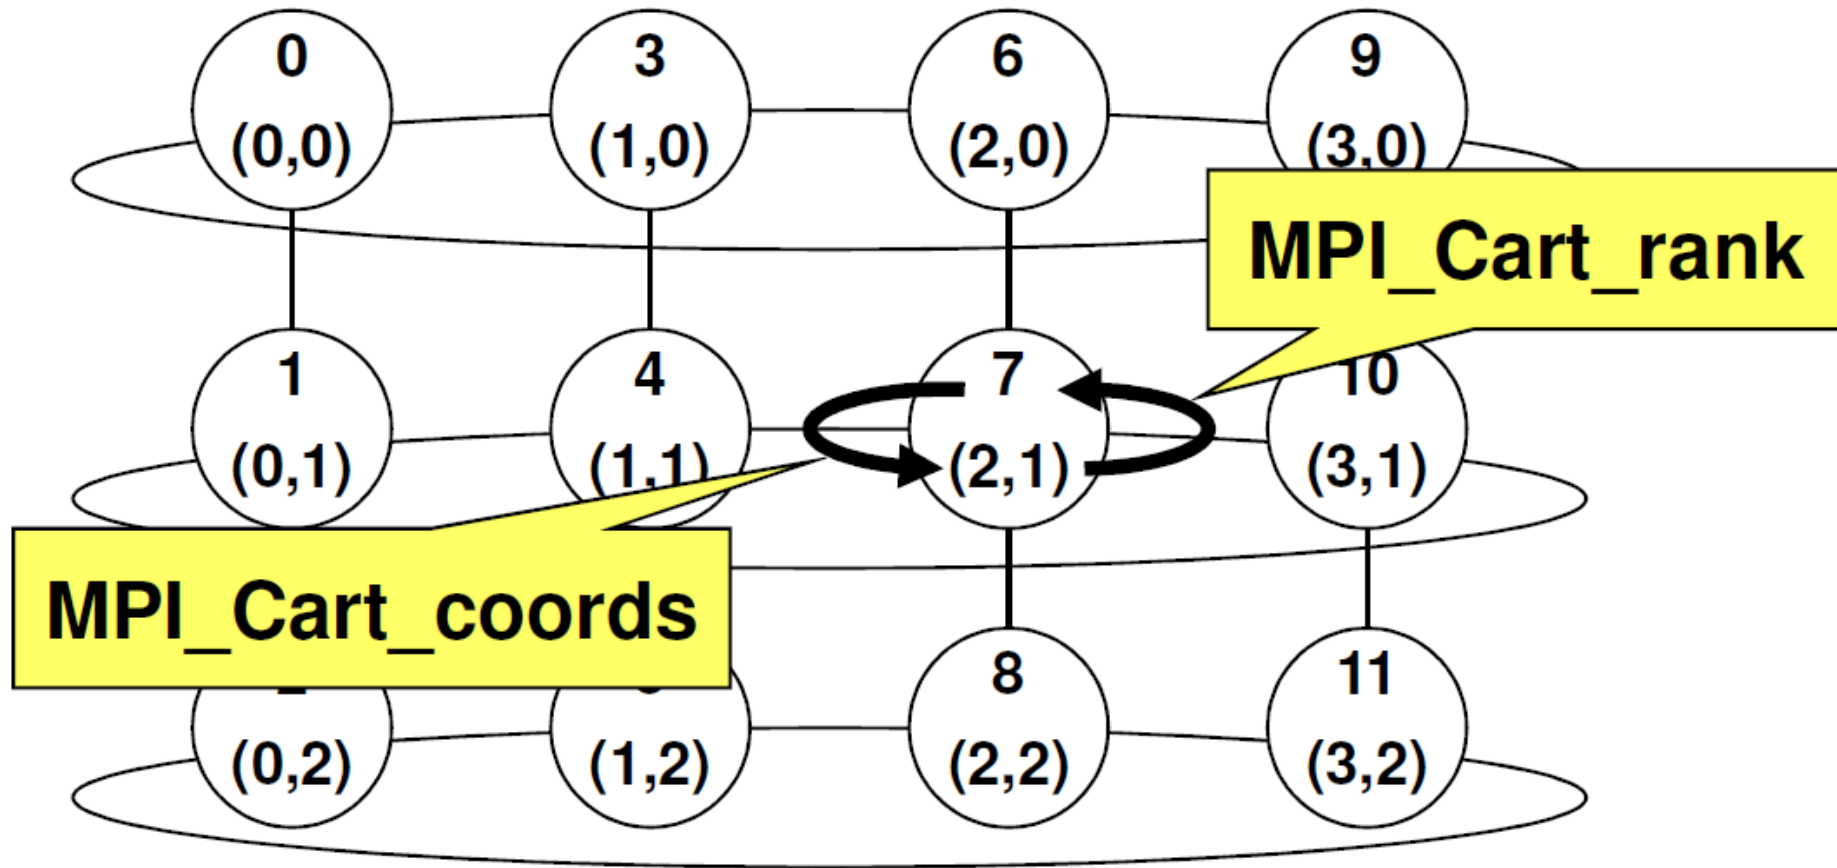
```
int MPI_Cart_coords(MPI_Comm comm_cart, int rank, int maxdims, int *coords)
```

- Mapping process grid coordinates to ranks



```
int MPI_Cart_rank(MPI_Comm comm_cart, int *coords, int *rank)
```

Cartesian Mapping Functions




Cartesian Mapping Functions

- Compute ranks of neighboring processes
- Returns `MPI_PROC_NULL` if there is no neighbor.
- In general: `MPI_PROC_NULL` can be used as source or destination rank in each communication. →
This communication will be a no-operation!

```
int MPI_Cart_shift(MPI_Comm comm_cart, int direction, int disp,  
                  int *rank_source, int *rank_dest)
```


Cartesian Mapping Functions



invisible input argument: **my_rank** in cart

`MPI_Cart_shift(cart, direction, displace, rank_source, rank_dest, ierror)`

example on

process rank=7

0 or
1

+1
+1


4
6

10
8

Cartesian Partitioning

- Cut a grid up into slices.
- A new communicator is produced for each slice.
- Each slice can then perform its own collective communications.

```
int MPI_Cart_sub( MPI_Comm comm_cart, int *remain_dims,  
                 MPI_Comm *comm_slice)
```




Outline

- Distributed Parallel Computing II
 - Non-Blocking Communication
 - Collective Communication
 - Derived Datatypes
 - Virtual Topologies
 - **Hybrid MPI+OpenMP**
- Quiz



Hybrid Parallel Programming

- Hybrid parallel programming:
Mix different parallel programming approaches to efficiently use available computing platforms (heterogeneous computing resources)
- Typical example:
Programming of clusters of shared memory nodes



Hybrid Parallel Programming


- Hierarchical system layout
- Hybrid programming seems *natural*
 - MPI between the nodes
 - Shared memory programming inside of each SMP node
 - OpenMP
 - MPI-3 shared-memory programming: not covered here
 - Accelerator programming (CUDA, OpenACC, etc): not covered here
 - And others.




Hybrid Parallel Programming

- Which programming model is best?


- “Pure” MPI?
(a.k.a. 1 MPI process per CPU core,
no shared-memory at all)



- “Fully hybrid”?
(a.k.a. 1 MPI process per node,
shared-memory within nodes)

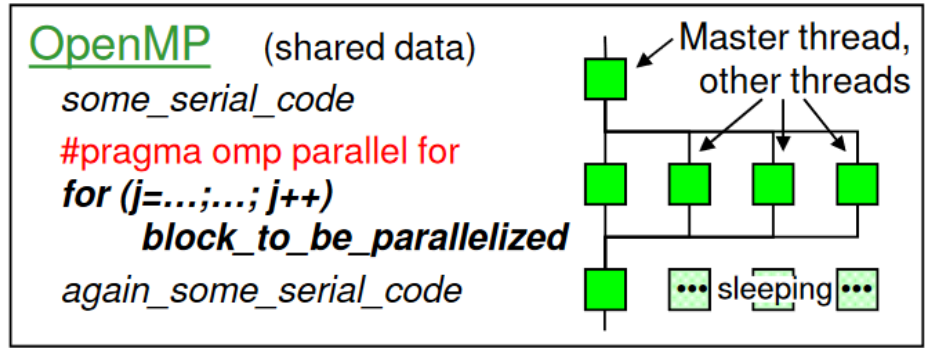
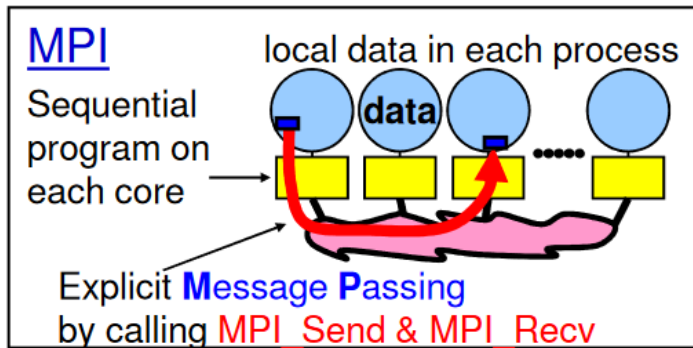



- “Mixed hybrid”?
(a.k.a. 1 MPI process per, e.g., CPU,
shared-memory within the CPUs;
other combinations possible, e.g., (cc)NUMA specific)



Problem specific ...

Typical Case: MPI+OpenMP



Hybrid MPI+OpenMP *Masteronly*

- **Masteronly:**
MPI only outside parallel OpenMP regions →
only the master thread issues MPI calls

```
#pragma omp parallel  
    numerical code  
/*end omp parallel */
```

```
/* on master thread only */  
MPI_Send (original data to halo areas in other SMP nodes)  
MPI_Recv (halo data from the neighbors)
```


Hybrid MPI+OpenMP *Masteronly*

- **Major Advantages**

- No message passing inside of the SMP nodes
- No topology problem

- **Major Disadvantages**

- All other threads are sleeping while master thread communicates
- All communicated data passes through the cache where the master thread is executing
- Strictly speaking: MPI-library must have been compiled with threading support (configure OpenMPI library build with: --enable-mpi-threads)
and must have been initialized with threading support
→ But masteronly approach will likely work without as well.

MPI Threading Support Handles

- **MPI_THREAD_SINGLE:** Only one thread will execute (→ similar to calling MPI_Init ())
- **MPI_THREAD_FUNNELED:** Only master thread will make MPI-calls
- **MPI_THREAD_SERIALIZED:** Multiple threads may make MPI-calls, but only one at a time (not covered here)
- **MPI_THREAD_MULTIPLE:** Multiple threads may call MPI, with no restrictions (not covered here)

```
int MPI_Init_thread( int * argc, char ** argv[],  
                    int thread_level_required,  
                    int * thread_level_provided);
```

Outline

- Distributed Parallel Computing II
 - Non-Blocking Communication
 - Collective Communication
 - Derived Datatypes
 - Virtual Topologies
 - Hybrid MPI+OpenMP
- **Quiz**

Quiz

- **Q1: How many Bytes would be required to store an MPI derived datatype based on a struct of two characters and one double?**
- **Q2: Can a non-blocking receive be combined with a blocking send?**
- **Q3: Give an example setup for defining a Cartesian MPI communication topology for a three-dimensional setup if you could use up to 1000 MPI processes (make assumptions for all other properties).**
- **Q4: What is the order of the forward/backward Euler method?**
- **Q5: Are Runge-Kutta methods single-step or multi-step methods?**