

Numerical Simulation and Scientific Computing I

Lecture 1: Introduction and Computer Architectures



Xaver Klemenschits, Paul Manstetten, and
Josef Weinbub



Institute for Microelectronics
TU Wien

nssc@iue.tuwien.ac.at

Outline


- **Introduction to the Lecture**
 - Course Goals
 - Rules
 - Survey
- **Computer Architectures**
 - Defining Computer Architecture
 - Memory Hierarchies
 - Instruction-Level Parallelism
 - Modern Multiprocessors
- **Quiz**

Team Presentation

- Josef Weinbub



- Paul Manstetten



- Xaver Klemenschits



nssc@iue.tuwien.ac.at

TISS

- Exam registration



TUWEL

- Lecture slides
 - Exercise handouts
 - General information and announcements
 - Student communication: Forum and Chat
- <https://tuwel.tuwien.ac.at/course/view.php?idnumber=360242-2022W>



The screenshot shows the TISS (TU Wien Information System) interface. The main content area displays course information for '360.242 Numerical Simulation and Scientific Computing I' for the 2022W semester. The course is listed as '2022W, VU, 3.0h, 6 ECTS'. A status bar indicates 'Erfolgreich aktualisiert'. Below this, there are tabs for 'Ankündigung', 'Durchführung', 'Kommunikation', and 'Terminkoordination'. The 'Durchführung' tab is active, showing a table of lecturers and their status.

Vortragende	Titel	Art	Tarif	Beauftragung	Abhaltung	Institut
Wieshub, Josef, Associate Prof. Dipl.-Ing. Dr.techn. BSc	1.50	Assoc.	Beauftragte	✓		E360-01
Manstetter, Paul, Univ.-Ass. Dipl.-Ing. (FH) Dr.techn. MSc	0.75	PostDoc	50%			E360-01
Klemensichits, Xaver, Projektleiter Dr.techn. MSc	0.75	Proj.Ass.	50%	✓		E360-01

Below the table, there is a 'Status Beauftragung' section with a table showing the status of the course. The table has columns for 'Offen', 'Beauftragte', and 'Abgelehnt'. The 'Summe Vortragende' row shows 0.00 for 'Offen', 3.00 for 'Beauftragte', and 0.00 for 'Abgelehnt'. The 'Summe StudienTutor_innen' row shows 0.00 for 'Offen', 0.00 for 'Beauftragte', and 0.00 for 'Abgelehnt'. The 'Gesamtsumme' row shows 0.00 for 'Offen', 3.00 for 'Beauftragte', and 0.00 for 'Abgelehnt'.

There is also a 'Mitwirkende' section with a table showing the status of the course. The table has columns for 'Stunden', 'Beauftragungstatus', and 'Institut'. The 'Mitwirkende' row shows 'No records found'.

At the bottom, there is a 'Vorschau' section with a list of links: 'Zum TUNEL Online-Kurs' and 'Zum LVA-Forum'. There is also a 'Merkmal' section with a list of features: 'Semesterwochenstunden: 3.0', 'ECTS: 6.0', 'Typ: VU Vorlesung mit Übung', and 'Format der Abhaltung: Präsenz'.

The TUWEL logo is circled in red in the bottom right corner of the screenshot.

General Goals

- Short introduction to many aspects of CSE
- Most topics will be treated in-depth in specialized courses
- Lots of hands-on work!
- Learning the basic vocabulary

Learning Outcomes

You will ...

- be able to select and apply fundamental methods of scientific computing;
- be able to judge the challenges regarding computing time and implementation effort;
- gain skills to solve inter-disciplinary, compute-intensive problems;
- be able to evaluate and analyze computational approaches;
- be able to scientifically formulate and extensively analyze compute-intensive problems as well as develop appropriate approaches.

What We Will Not Cover

- GPUs →
360.252
VU Computational Science on Many-Core Architectures
Winter term 2020
- Distributed Parallel Computing, Classification and Analyses
of Partial Differential Equations, Finite Volume Method, Finite
Element Method, Molecular Dynamics, Fluid Dynamics →
360.243
Numerical Simulation and Scientific Computing II
Summer term 2021

Course Calendar - Lectures

Lecture: 14:00 – 15:30

Date	Subject
October 6	Introduction and Computer Architectures
October 13	Serial Code Optimization
October 20	C++ Optimization, Intro Finite Difference Method
October 27	Finite Difference Method and Finite Precision Floating Point Arithmetic
November 3	Numerical Linear Algebra I
November 10	Numerical Linear Algebra II
November 17	Random Numbers and Monte Carlo Integration
November 24	Shared Memory Parallel Computing
December 1	Algorithm Analysis and Data Structures
December 15	Mesh Generation
December 22	Build Systems and Debugging
January 12	Version Control and Open Source Licenses
January 26	Written Exam

Rules - Exercises

- Three mandatory group exercises over the whole lecture
- Group size: 4 – *Find your group* via TUWEL forum/chat:
 - Send your group info until October 12 to nssc@iue.tuwien.ac.at
- Copying code not allowed (**we also check previous years**):
If identified – **0 points for all groups** involved!
- Each exercise will be graded separately from 0 to 10 points
- Exercise threshold for final exam: **sum of ex. grades ≥ 21.0**
- **Submission deadlines are hard deadlines**
- Submit all files as a ZIP file on TUWEL:
 - Source code, I/O, plots and a Makefile (**no binary files!**)
 - Also, a plain .txt file should there be any written answers
- You will be provided a guide to setup a development environment: VirtualBox virtual machine running Linux Mint
- *Anyone here who did the exercises in the past?*

Course Calendar - Exercises

Exercise Meetings: Right after the lecture

Send your group info until October 12 to nssc@iue.tuwien.ac.at

Date	Exercise 1
October 13	Handout session
October 20	Support session I
October 27	Support session II
November 3	Submission deadline
Date	Exercise 2
November 3	Handout session
November 10	Support session I
November 17	Support session II
November 24	Submission deadline
Date	Exercise 3
November 24	Handout session
December 1	Support session I
December 15	Support session II
December 22	Submission deadline

Rules - Quizzes

- In the end of each lecture, you will receive 5 questions
- 3 questions reviewing the current lecture
- 2 questions preparing for the next
- Discussion in the beginning of next class
- Participation is voluntary (does not count towards your final grade) but encouraged
- These questions will help you learn for the exam!

Rules - Exam

- The examination will be in written form
- Expect a tough 2–3-hour exam
- Mix of theory and practical questions:
 - Make sure that you are an expert in all exercises!
- Reserve appropriate time for preparations
- Learning material:
Lectures, exercises, and quizzes

Grading Scheme

- Your achievements in the exercises and the exam are weighted for the final grade:
 - Exam: $\frac{3}{4}$
 - Exercises: $\frac{1}{4}$
- But, don't forget: You will need to achieve a minimum of 21 points on the exercises to be eligible for the exam!

Computer Resources

- It is expected that you have access to a modern, x86-based, personal computer (laptop, home desktop, etc.)
- GNU-Linux operating system is required;
a guide to setup a ready-to-use virtual machine will be provided

Programming Languages

- We assume you have experience with at least one of the following:
 - C, C++, Python, Fortran
- The programming exercises are to be handed in **C++**
 - The **C++11** standard is preferred but not required
- For data analysis and visualization, use **Python**
 - Including **numpy** and **matplotlib**
- More information in the respective exercise handouts
- Resources:
 - https://catalogplus.tuwien.ac.at:443/UTW:UTW:UTW_alma5167982990003336 (Book, A Tour of C++)
 - <https://cimec.org.ar/~mstorti/C/CURSOS/advc++3.pdf> (Presentation, good if you already know C)

Class Survey

- Background:
 - Engineering (Mechanical/Electrical/Civil/Materials), CS, Math, Physics, Chemistry, Life Sciences, Other
- Known Programming Languages:
 - C, C++, Python, Fortran, MATLAB, Julia, Java, JavaScript, Ruby, R, BASIC, COBOL, Others
- Differential Equation Methods:
 - Euler, Runge-Kutta, Finite Difference, Finite Volume, Finite Element
- Linear Algebra:
 - BLAS/LAPACK, Sparse Matrices, Other Frameworks
- Parallel Programming:
 - pthreads, OpenMP, HPF, MPI, CUDA, Other
- Software Engineering:
 - Makefile, Cmake, Git, GDB

Outline


- Introduction to the Lecture
 - Course Goals
 - Rules
 - Survey
- **Computer Architecture**
 - Defining Computer Architecture
 - Memory Hierarchies
 - Instruction-Level Parallelism
 - Modern Multiprocessors
- Quiz

Reference Book

- Computer Architecture: A Quantitative Approach 6th Edition
 - Authors: John L. Hennessy, David A. Patterson
 - A few physical copies on the main library:
https://catalogplus.tuwien.ac.at:443/UTW:UTW:UTW_alma2172502750003336
- The older 5th edition is fully available online from the library:
 - https://catalogplus.tuwien.ac.at:443/UTW:UTW:UTW_alma5172773450003336
- Chapters 1, 2 and 3

42 Years of Microprocessor Trend Data

42 Years of Microprocessor Trend Data



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2017 by K. Rupp

Computer Categories

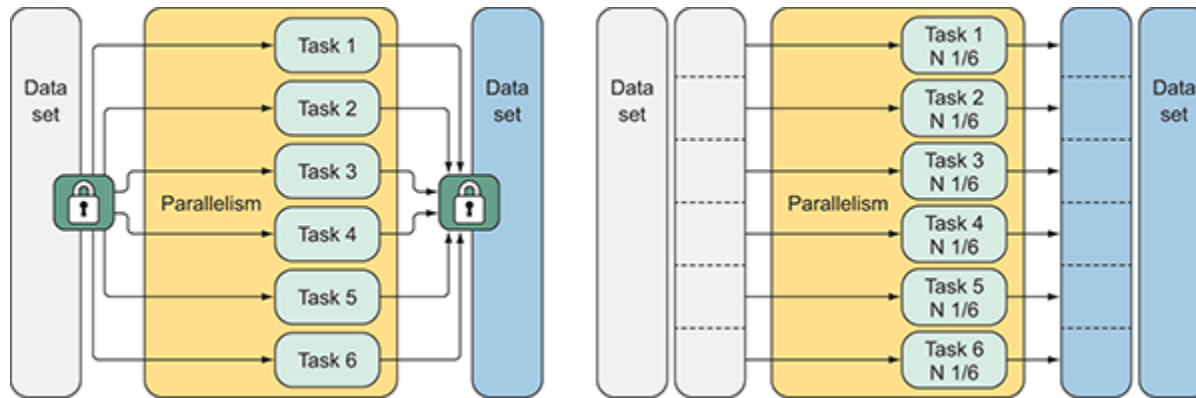
- Embedded Computers/Internet of Things (IoT)
 - Computers inside everyday things! Cars, microwaves, dishwashers...
 - IoT = embedded computer + wireless internet
 - Cheap!
- Personal Mobile Devices
 - Probably we are all carrying one right now!
 - Energy-efficient
- Desktop Computing
 - Laptops, desktops...
 - Price/performance is the key!
- Servers
 - Availability, scalability, throughput
- Clusters/Warehouse-Scale Computers
 - A networked collection of servers or desktop computers
 - *Where the cloud lives*
 - Subclass: Supercomputers

Classes of Parallelism

- Parallelism in Applications:

Task-level parallelism (TLP)


Data-level parallelism (DLP)




- Architectural parallelism:

- Instruction-level parallelism (e.g., pipelining) – *will come back later*
- Vector architectures (e.g., GPUs) – *not covered*
- Thread-level parallelism (i.e., multi-threading) – *will come back later*
- Request-level parallelism (e.g., internet services) – *not covered*

Serial Problem



Parallel Problem



Flynn's Taxonomy

- Single instruction stream, single data stream (SISD)
 - Classical uniprocessor
 - Instruction-level parallelism
- Single instruction stream, multiple data streams (SIMD)
 - Vector architectures/GPUs
 - Multimedia/Vector extensions (SSE, AVX...)
 - Data-level parallelism
- Multiple instruction streams, single data streams (MISD)
 - Task replication for fault detection
 - Very uncommon
- Multiple instruction streams, multiple data streams (MIMD)
 - Each processor fetches and operates on own data: task-level parallelism
 - Tightly-coupled MIMD (e.g. multicore CPU): thread-level parallelism
 - Loosely-coupled MIMD (e.g. cluster): request-level parallelism

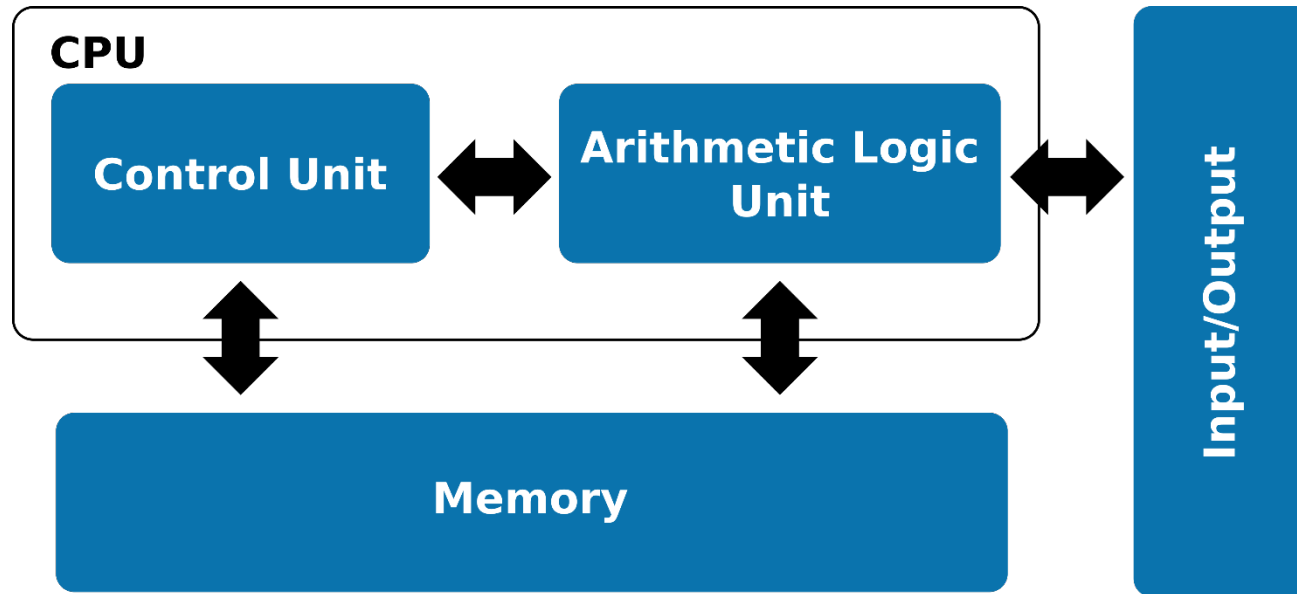
Defining Computer Architecture

- “Old view” of computer architecture:
 - Instruction Set Architecture (ISA) design:
register, memory addressing, operands, control flow instructions, encoding, etc.
- “Today’s view” of computer architecture:
 - Not only ISA but also microarchitecture and specific hardware
 - Everything we discussed put together!

Processor

- Processor
 - Integrated electronic circuit
 - Performs operations on data
 - Many processors in a computer
 - CPU: Central processing unit
 - GPU: Graphics processing unit
 - TPU: Tensor processing unit
 - DSP: Digital signal processors

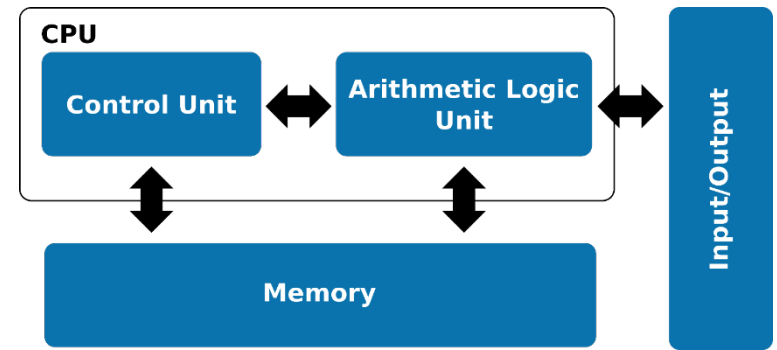
CPU's von Neumann Architecture



- Program feeds control unit
- Stored in memory together with data
- Data required by arithmetic unit

Von Neumann Architecture - Issues

- Memory bottleneck
 - Instructions and data must be fed to control and arithmetic units
 - Speed of memory interfaces limits computational performance:
Von Neumann bottleneck
- Inherently sequential architecture
 - Example of SISD
 - Still allows for interesting tricks using instruction-level parallelism...
- However...
 - It is the only game in town.
 - We better make the best out of it!




Instruction Set Architecture (ISA) - Examples

- x86
 - Desktop computing, servers, clusters
- ARMv8
 - Mostly used for personal mobile devices (both Android & iOS)
- IBM Power9
 - Common in high-end supercomputers
- Other examples
 - SPARC, RISC-V, PowerPC, Itanium, Cell, Sunway...
- Code compiled for one ISA will not work on another
 - Also: program outputs might not be read by different ISAs (endianness)

Bandwidth and Latency

- Bandwidth (also: throughput)
 - Total work done in a certain time
- Latency (also: response time)
 - Time between start and completion of an event
- Example: loading a truck with hard drives and driving to another location
 - Very high bandwidth
 - Very high latency
- Bandwidth improves faster than latency!




Principle of Locality

- We observe in practice a kind of *Pareto principle*
 - A program spends 90% of its time on 10% of the code
- Programs tend to reuse data and instructions used recently
- Temporal locality: recently accessed items are likely to be needed again
- Spatial locality: items with near addresses tend to be referenced close in time

Memory Hierarchy - Basics

- Ideally, we would like unlimited memory with zero latency
- Fast memory is much more expensive per bit than slower memory
- Solution: organize memory system into a hierarchy
 - Entire addressable memory available in largest, slower memory
 - Incrementally faster but smaller memories, containing a subset of the memory below
- Temporal and spatial locality enables that many references are found in smaller/faster memories

Memory Hierarchy




Memory Hierarchy - Latency

Event	Real time	“Normalized” time
1 CPU cycle	0.3ns	1s
Level 1 cache access	0.9ns	3s
Level 2 cache access	2.8ns	9s
Level 3 cache access	12.9ns	43s
Main memory access	120ns	6min
Solid-state disk I/O	50-150μs	2-6 days
Rotational disk I/O	1-10ms	1-12 months
Internet: SF to NYC	40ms	4 years
Internet: SF to UK	81ms	8 years
Internet: SF to Australia	183ms	19 years
OS virtualization reboot	4s	423 years
SCSI command time-out	30s	3000 years
Hardware virtualization reboot	40s	4000 years
Physical system reboot	5min	32 millennia

Source: Gregg: Systems Performance, Prentice Hall, 2013

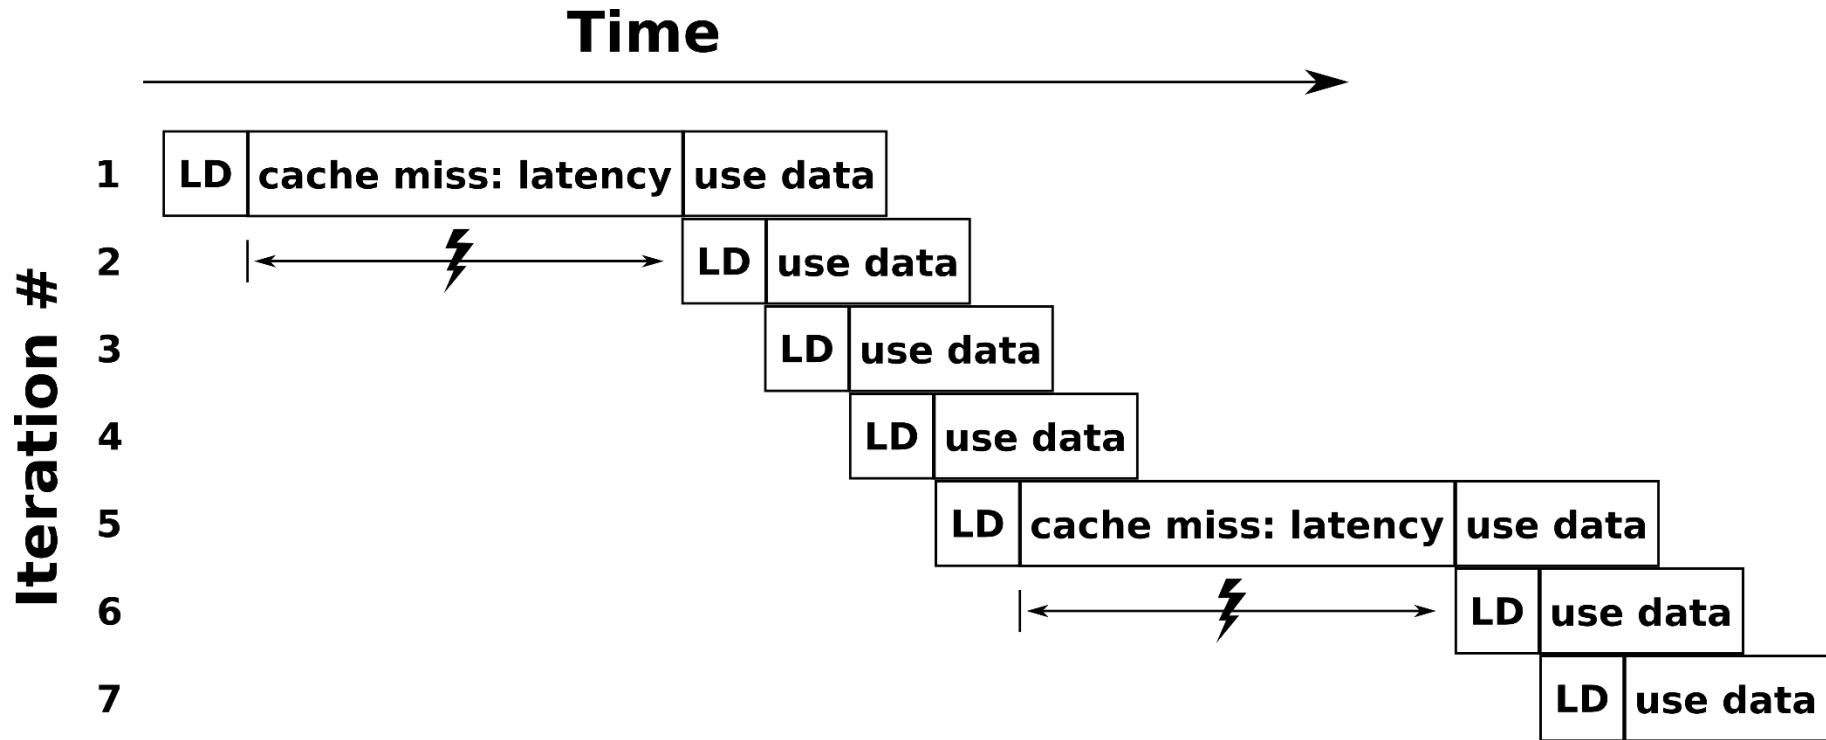
Memory Hierarchy – Memory Gap




Memory Hierarchy - Caches

- Caches are low-capacity ($>10\text{MB}$), high-speed memory integrated on CPU die
 - Usually 2-3 levels: L1-L2-L3
 - L1 usually split into instructions (L1I) and data (L1D)
 - L2 usually unified
 - L3 usually shared between all cores
- Data load request:
 - If available: cache hit
 - If not available: cache miss
- Cache miss: fetch data from lower levels in hierarchy
 - Also fetch other data within the block (spatial locality)

Memory Hierarchy - Caches



Memory Hierarchy - Caches

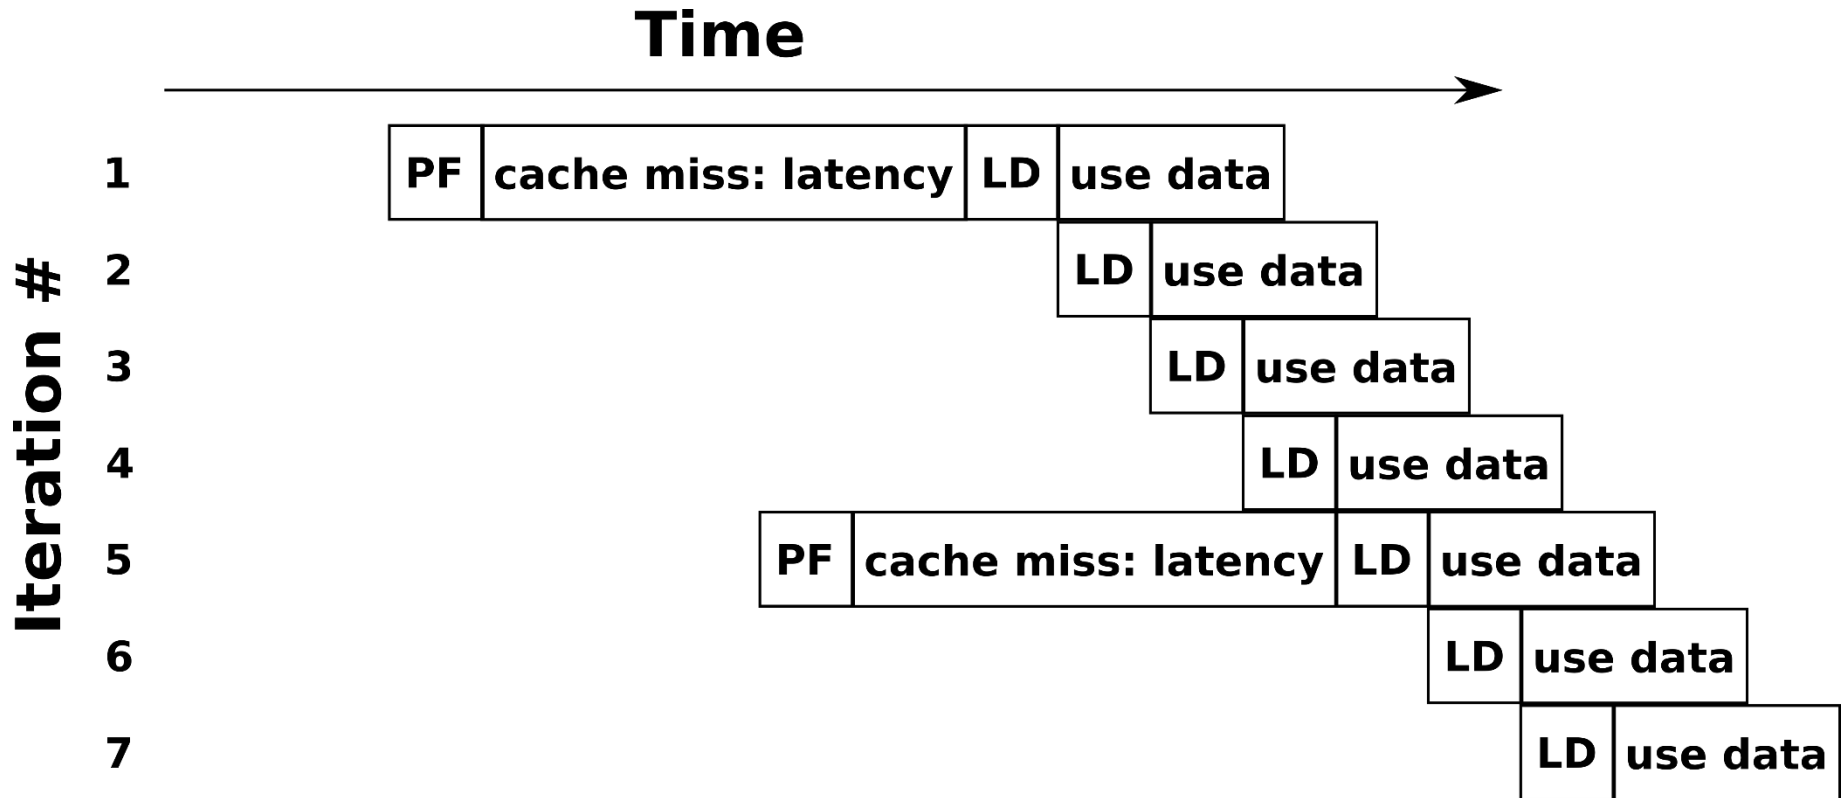


- Shared caches in a multicore processor (e.g. Intel Nehalem)

Memory Hierarchy - Prefetch

- Provides caches with data ahead of time
 - Hides latency
 - Asynchronous memory operation
- Software prefetching by compilers
 - Interleaves special prefetching instructions in running pipeline
 - “Touches” certain cache lines in advance
- Hardware prefetchers: identify patterns to read ahead
- Both hardware and software prefetchers are used today
 - These techniques can only do so much
 - Significant burden on memory subsystem
 - Tip: provide long continuous data streams

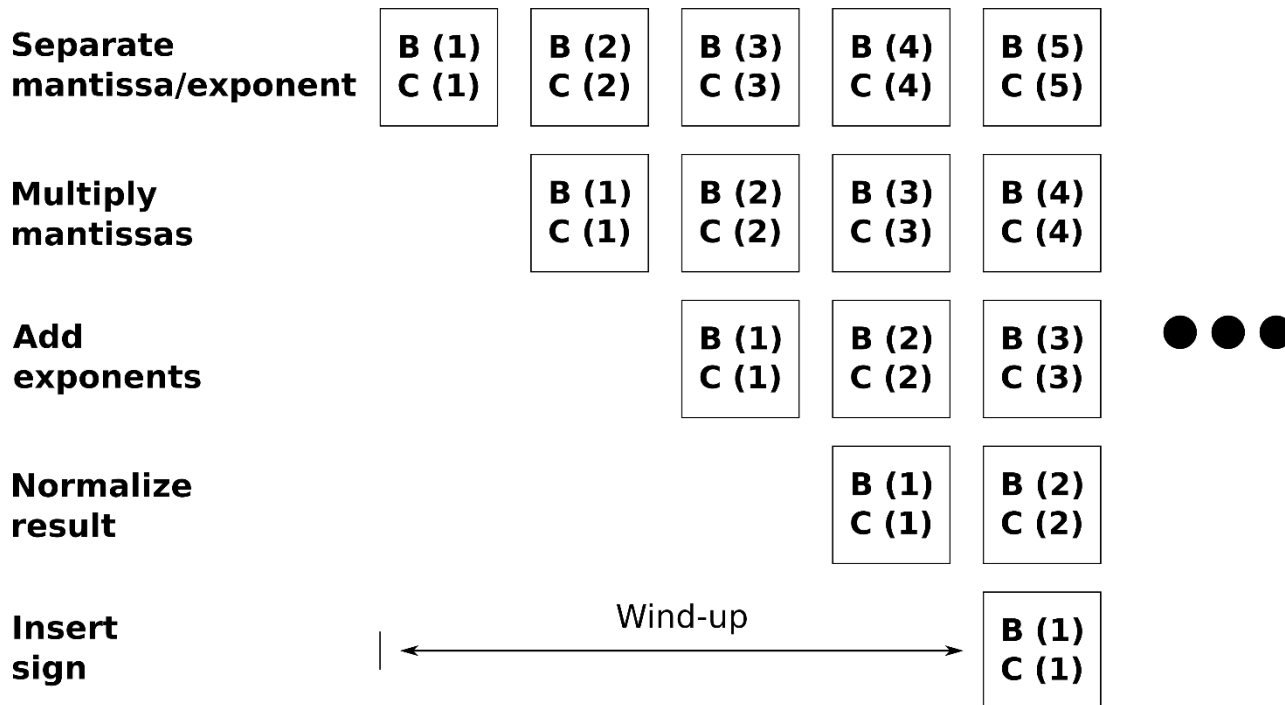
Memory Hierarchy - Prefetch



Instruction-Level Parallelism (ILP)

- Exploitation of data-level parallelism at modest levels
- Key idea: modern CPUs can execute more than one instruction simultaneously
- Main concepts:
 - Pipelining
 - Loop unrolling
 - Branch prediction

Pipelining



- Simplified example: vector product $A(:) = B(:) * C(:)$
- Complex operations require more than one cycle
- Analogous to an assembly line
 - Each “worker” does a subtask, but they can all work simultaneously if the line is full

Pipelining

- The classical example of ILP
 - Processors can handle this since the 1980's!
- However, many known issues:
 - Inefficient for short and tight loops
 - Operations with long latency (e.g., square root)
 - Pipeline *bubble* when low degree of pipelining
 - Pipeline *stall* if data is not loaded in time

Other ILP Techniques

- Loop Unrolling
 - Basic loop transformation technique
 - Replicates the loop body multiple times, adjusting the termination code
 - Requires independent loop iterations
 - Helps filling up the pipeline faster!
- Branch Prediction
 - E.g., “if-else-then”
 - In practice, implemented in-hardware (dynamic branch prediction)
 - It is too costly to wait for a decision regarding which branch
 - So, guess a branch!
 - If a branch has been taken previously, simply assume it will be taken again
 - Security issues: Meltdown and Spectre


Modern Processors - Examples

- Consumer-level
 - Intel i3-i9, AMD Ryzen 3-9, Apple Mx
- Smartphone SOCs
 - Qualcomm Snapdragon, Samsung Exynos, Huawei Kirin, Apple Ax
- Cluster-tailored hardware
 - IBM Power9, Intel Xeon, AMD EPYC

Increase Processor Core Count

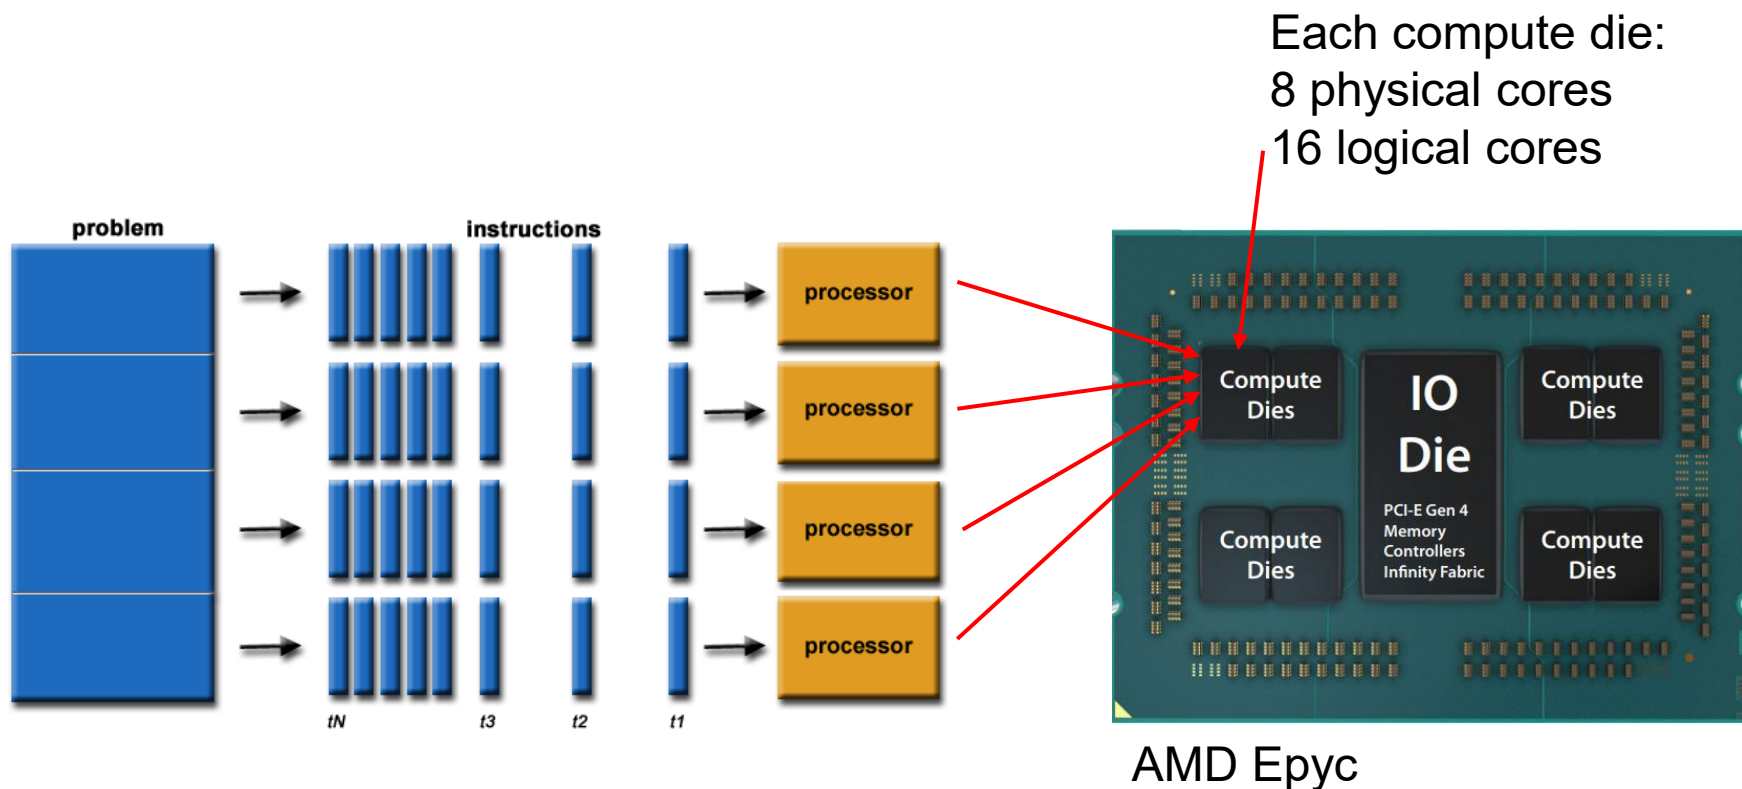
Two ways to increase the core count of CPUs:

- Multicore: Add additional cores with complete ISA
 - aka “physical cores”
- Simultaneous Multithreading: Duplicate registers and control; share cache and pipeline
 - aka “logical cores” (= 2x physical cores)



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2017 by K. Rupp

Why Increase Core Count (Again)?

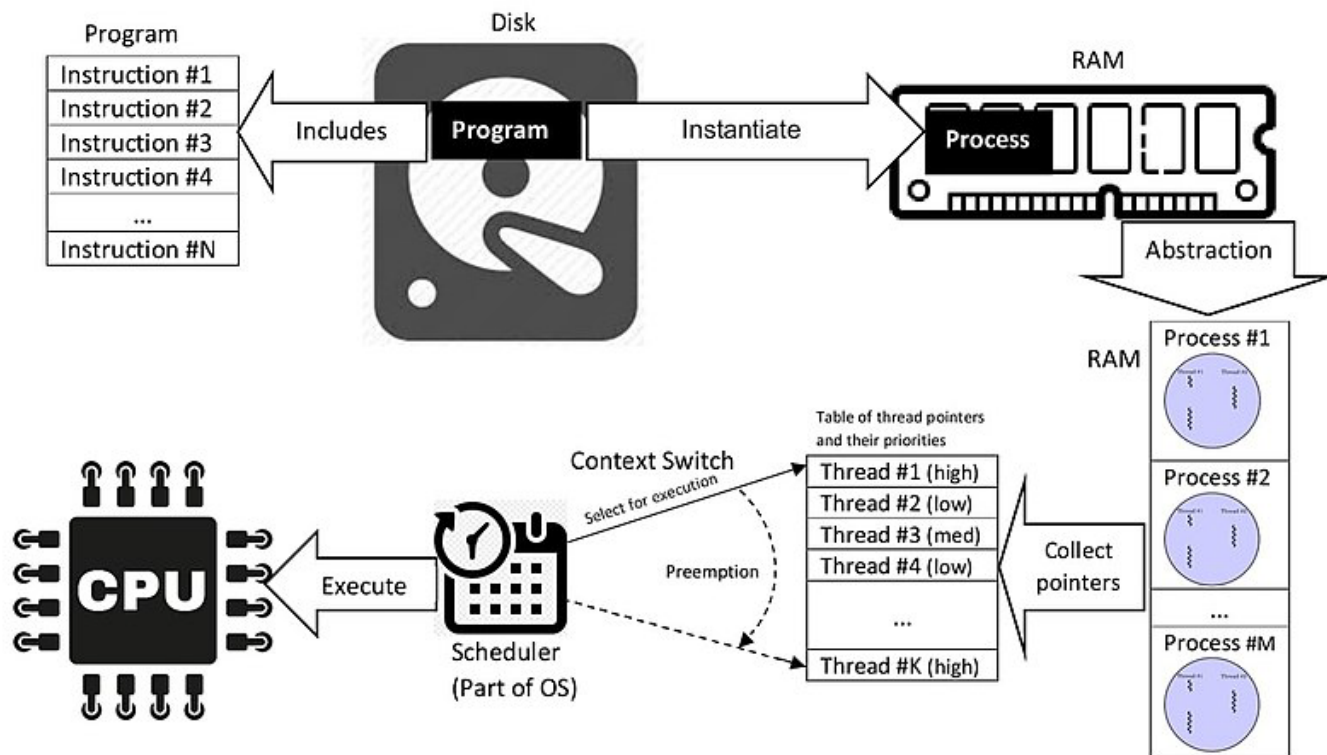


“Problem” can be:

- A single parallel program: single process, multiple threads
- Several serial programs: single process, single thread

Program vs Process vs Thread

- Program: Collection of instructions
- Process: Execution of a program
- Thread: Execution of smallest sequence of instructions
 - A single process can contain multiple threads and share resources




Multithreading

- Temporal Multithreading

- A CPU can switch between multiple threads very quickly
- So fast that users might not realize!
- Why even old computers don't stop after you click on something...
- However, only one thread is executed

- Simultaneous Multithreading (SMT)


- Introduced by Intel back in the Pentium 4 era (Hyperthreading)
- CPU appears to be composed of more cores (aka “logical” cores)
- This is achieved by partly repeating the architecture (i.e. registers, stack, instruction pointers)
- More than one thread is executed but also requires more overhead



CPU Specifications


Total Cores	4
Total Threads	8

No Simultaneous Multithreading



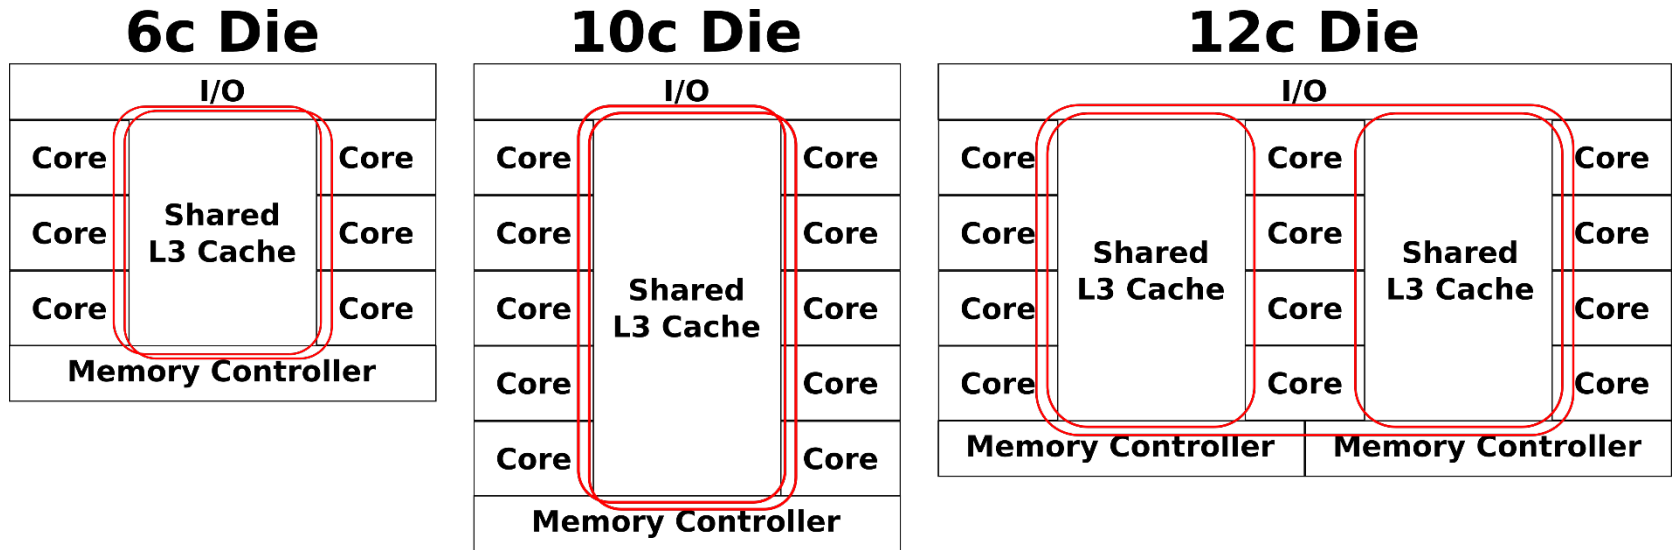
- White blocks in pipelines denote bubbles

Simultaneous Multithreading



- Two-way simultaneous multithreading
- Two threads share caches and pipelines but retain respective architectural state

Multicore Processors




- Example: Intel Ivy Bridge EP
- Solution to the problem of single-core performance...
 - Use the transistor budget to **add more cores!**
- Usually **shares L3 cache** (at least partly...)
- Does not solve the memory bandwidth problem!

Cache Coherence

- Cache coherence is required in all cache-based multi-processor systems
- Copies of cache lines can reside in different caches
- Imagine: One cache line gets modified ..
- Cache coherence protocols ensure consistent view of memory at all times
- Coherence traffic can hurt application performance if the same cache line is frequently modified by different processors (i.e. false sharing)
- Implemented in the hardware (CPU or chipsets)

Cache Coherence



MESI

- M: modified
- E: exclusive
- S: shared
- I: invalid


(there are others as well, e.g., MOESI, MESIF)

Example: P1, P2 modify the two parts A1, A2 of the same cache line (CL) in caches C1 and C2. The MESI coherence protocol ensures consistency between cache and memory:

1. C1 requests exclusive CL ownership
2. Set CL in C2 to state I
3. CL has state E in C1 → modify A1 in C1 and set to state M
4. C2 requests exclusive CL ownership
5. Evict CL from C1 and set to state I
6. Load CL to C2 and set to state E
7. Modify A2 in C2 and set to state M in C2

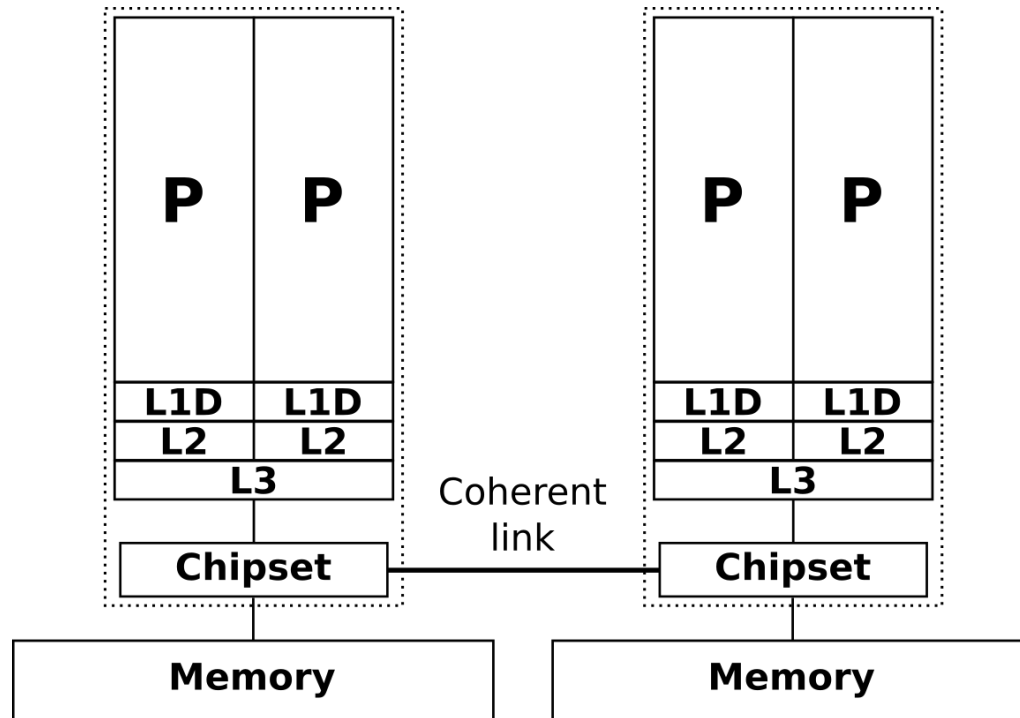
Question: Can we assume that access (latency and bandwidth) to memory is the same for all threads executed on a processor?

Uniform Memory Access (UMA)




- Two examples where access (latency and bandwidth) are the same for all threads: Both offer UMA!
- Latency and bandwidth same for processors/memory
- Also known as: Symmetric multiprocessing (SMP)
- Limited scalability: memory bus contention
- E.g. typical single multi-core processor chips are UMA

Cache Coherent Non-Uniform Memory Access (ccNUMA)



- Memory is physically distributed but logically shared
- Better scalability:
Small number of CPUs compete for shared memory bus
- Physical layout similar to distributed-memory case

Dual-Socket Setup



Outline

- Introduction to the Lecture
 - Course Goals
 - Rules
 - Survey
- Computer Architectures
 - Defining Computer Architecture
 - Memory Hierarchies
 - Instruction-Level Parallelism
 - Modern Multiprocessors
- **Quiz**

Quiz – Question 1

- What is the minimum sum of exercise points you need to be eligible for the exam?

Quiz – Question 2

- Read the specification of this CPU

https://en.wikichip.org/wiki/intel/core_i9/i9-13900k

and classify it according to:

- Computer category
- Supported forms of architectural parallelism
- Flynn's taxonomy
- Instruction set architecture
- Memory hierarchy (cache levels)
- Type of multithreading
- Type of multicore architecture

Quiz – Question 3

- Caches are very expensive. Why do we need them at all?

Quiz – Question 4

- What are simple ways you can use a compiler to make your program run faster?

Quiz – Question 5

- Why are SIMD extensions (e.g. AVX-512) not considered instruction-level parallelism?