

Numerical Simulation and Scientific Computing I

Lecture 7: Random Numbers & Monte Carlo Integration



Xaver Klemenschits, Paul Manstetten, and
Josef Weinbub



Institute for Microelectronics
TU Wien

nssc@iue.tuwien.ac.at

Quiz

- Q1: Represent the following matrix in the CRS format

$$\begin{bmatrix} 1 & 4.6 & 0 & 0 & 0 \\ 0 & 0 & 8.5 & 3.7 & 0 \\ 0 & 6 & 2.7 & 0 & 0 \\ 0 & 4.6 & 0 & 4.8 & 9.4 \\ 0 & 0 & 5.6 & 0 & 1 \end{bmatrix}$$

Quiz

- Q1: Represent the following matrix in the CRS format

$$\begin{bmatrix} 1 & 4.6 & 0 & 0 & 0 \\ 0 & 0 & 8.5 & 3.7 & 0 \\ 0 & 6 & 2.7 & 0 & 0 \\ 0 & 4.6 & 0 & 4.8 & 9.4 \\ 0 & 0 & 5.6 & 0 & 1 \end{bmatrix}$$

- 1st step: Put the values in row-major order

$$V = [1 \quad 4.6 \quad 8.5 \quad 3.7 \quad 6 \quad 2.7 \quad 4.6 \quad 4.8 \quad 9.4 \quad 5.6 \quad 1]^T$$

Quiz

- Q1: Represent the following matrix in the CRS format

$$\begin{bmatrix} 1 & 4.6 & 0 & 0 & 0 \\ 0 & 0 & 8.5 & 3.7 & 0 \\ 0 & 6 & 2.7 & 0 & 0 \\ 0 & 4.6 & 0 & 4.8 & 9.4 \\ 0 & 0 & 5.6 & 0 & 1 \end{bmatrix}$$

- 2nd step: Assign the corresponding column indices

$$V = [1 \quad 4.6 \quad 8.5 \quad 3.7 \quad 6 \quad 2.7 \quad 4.6 \quad 4.8 \quad 9.4 \quad 5.6 \quad 1]^T$$

$$JA = [0 \quad 1 \quad 2 \quad 3 \quad 1 \quad 2 \quad 1 \quad 3 \quad 4 \quad 2 \quad 4]^T$$

Quiz

- Q1: Represent the following matrix in the CRS format

$$\begin{bmatrix} 1 & 4.6 & 0 & 0 & 0 \\ 0 & 0 & 8.5 & 3.7 & 0 \\ 0 & 6 & 2.7 & 0 & 0 \\ 0 & 4.6 & 0 & 4.8 & 9.4 \\ 0 & 0 & 5.6 & 0 & 1 \end{bmatrix}$$

- 3rd step: Identify the indices in V where a new row starts

$$V = [1 \quad 4.6 \quad 8.5 \quad 3.7 \quad 6 \quad 2.7 \quad 4.6 \quad 4.8 \quad 9.4 \quad 5.6 \quad 1]^T$$

$$JA = [0 \quad 1 \quad 2 \quad 3 \quad 1 \quad 2 \quad 1 \quad 3 \quad 4 \quad 2 \quad 4]^T$$

$$IA = [0 \quad 2 \quad 4 \quad 6 \quad 9 \quad 11]^T$$

Quiz – Poll 1

- Q2: Under which conditions would you use the CG method instead of GMRES?
- A) The matrix is indefinite
- B) The problem benefits from preconditioning
- C) The problem is ill-posed
- D) All eigenvalues are positive
- E) The LU factorization is known

Quiz – Poll 1

- Q2: Under which conditions would you use the CG method instead of GMRES?
- A) The matrix is indefinite
- B) The problem benefits from preconditioning
- C) The problem is ill-posed
- **D) All eigenvalues are positive**
- E) The LU factorization is known

Quiz – Poll 2

- Q3: What information from a matrix A could be useful to help choosing an adequate solver?
- A) Maximum norm
- B) Set of eigenvalues
- C) QR decomposition
- D) Transpose

Quiz – Poll 2

- Q3: What information from a matrix A could be useful to help choosing an adequate solver?
- A) Maximum norm
- **B) Set of eigenvalues**
- C) QR decomposition
- D) Transpose

Outline

- Introduction to Randomness
- Random Number Generators (RNGs)
 - Linear Congruential Generators
 - Mersenne Twister
 - Hardware
- C++11 `<random>`
- Monte Carlo Integration

Take-home message

- Be careful with your choice of RNG
- Seeding is a consequence of pseudo-RNG
- Watch out for pitfalls on parallel systems

Main References

- Random Numbers and Computers
 - Author: Ronald T. Kneusel
 - eBook available:
https://catalogplus.tuwien.ac.at:443/UTW:UTW:TN_springer_s978-3-319-77697-2_453445

Additional References

- Handbook of Monte Carlo Methods
 - Authors: D. P. Kroese, T. Taimre, Z. I. Botev
 - eBook available:
https://catalogplus.tuwien.ac.at:443/UTW:UTW:TN_scopus2-s2.0-84949783693
- Monte Carlo Strategies in Scientific Computing
 - Author: Jun S. Liu
 - https://catalogplus.tuwien.ac.at:443/UTW:UTW:UTW_alma2146615720003336

Randomness

- How to formally define randomness?
- Philosophical conundrum: does randomness even exist?
 - Good thing that it is not our job to figure it out!
- Too hard - operational definition of a random sequence [Kneusel]:
 - a sequence of numbers n_i , within some bounded range, where it is not possible to predict n_{k+1} from any combination of preceding values $n_i, i = 0, 1, \dots, k$.

“True” Random vs. Pseudorandom

- “True” Randomness


- From here on, we will **assume that these exist**
- Found in the physical world
- Examples?

- Pseudorandomness

- [Kneusel]: A pseudorandom sequence is a **deterministically** generated sequence of numbers that is **indistinguishable** from a true random sequence of numbers
- Perfect indistinguishability is impossible – but how much is enough?
- Determinism: introduces the concept of a **seed**
 - **Seed**: starting point of the deterministic sequence

Uniform Random Numbers

- Even if n_{k+1} is truly non-predictable, it can still have a distribution – shows up in a histogram
 - Drawing from a **parent distribution**
- Uniform (real) Distribution: $n_k \in [0,1)$
 - Unless otherwise stated, it is **often the default**



Linear Congruential Generators

- Nomenclature

- a : multiplier
- c : increment
- m : modulus
- x_0 : seed

$$x_{n+1} = (ax_n + c) \bmod m$$

- Notes

- Generates integers $\in [0, m)$. Floats: $f_n = \frac{x_n}{m}$
- **Periodicity** is a big problem – can be lower than m
- If $x_n = x_m$ then $x_{n+1} = x_{m+1}$
- The method is defined by the choice of a , c and m

- Examples

- GCC: $a = 1103515245$, $c = 12345$, $m = 2^{31} - 1$
- C++11 minstd_rand: $a = 48271$, $c = 0$, $m = 2^{31} - 1$

LCGs – Poll 3

- What is the issue of using this generator with $m = 2^{31} - 1$ to generate a double?

$$x_{n+1} = (ax_n + c) \bmod m$$

- A) The periodicity is not high enough to cover the 53-bit significand
- B) It is impossible to seed this generator with a double
- C) You cannot generate a double from an integer
- D) The equation $f_n = \frac{x_n}{m}$ doesn't hold anymore

LCGs – Poll 3

- What is the issue of using this generator with $m = 2^{31} - 1$ to generate a double?

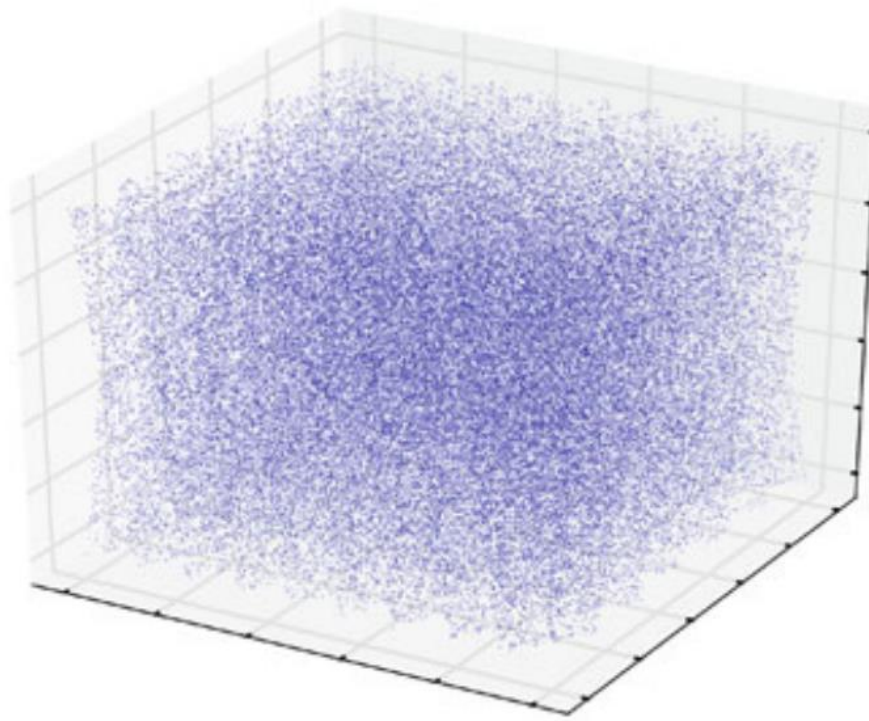
$$x_{n+1} = (ax_n + c) \bmod m$$

- **A) The periodicity is not high enough to cover the 53-bit significand**
- B) It is impossible to seed this generator with a double
- C) You cannot generate a double from an integer
- D) The equation $f_n = \frac{x_n}{m}$ doesn't hold anymore

A Lesson from History: RANDU

- How bad can it be if we choose poorly?
 - RANDU: $a = 65593$, $c = 0$, $m = 2^{31}$



$$x_{n+2} = 6x_{n+1} - 9x_n$$



A Lesson from History: RANDU


- How bad can it be if we choose poorly?
 - RANDU: $a = 65593$, $c = 0$, $m = 2^{31}$

$$x_{n+2} = 6x_{n+1} - 9x_n$$



Mersenne Twister

- MT is the current *de facto standard* for non-cryptographic applications



- Larger memory and construction requirements!

Source: Kneusel: Random Numbers and Computers, 1st ed., 2018, Springer

Choosing your RNG

- Performance is **only part of the equation**
- MINSTD is only 32 bits
 - a , c and m are linked
 - **Don't mess around with constants!**
- The performance penalty of MT can be mild (up to 50%)
- Other RNG may be available!
 - May involve other performance tradeoffs...
- C rand()
 - NOOOOOOOO ($RAND_MAX \geq 32767$)

Generator	Time (s)
MINSTD	9.164 ± 0.006
xorshift32	11.097 ± 0.004
xorshift128+	11.853 ± 0.028
CMWC	12.671 ± 0.003
Middle Weyl	12.849 ± 0.028
xorshift1024*	12.985 ± 0.004
Mersenne Twister	14.353 ± 0.007
KISS64	17.252 ± 0.015
Philox	18.097 ± 0.048
Threefry	32.442 ± 0.010

Source: Kneusel: Random Numbers and Computers, 1st ed., 2018, Springer

<https://channel9.msdn.com/Events/GoingNative/2013/rand-Considered-Harmful>

Choosing your RNG – Poll 4

- How does MT compare wrt. MINSTD?
 - Higher performance: lower time to generate 1 RN
- A) Higher performance, higher periodicity, higher memory requirements
- B) Lower performance, higher periodicity, higher memory requirements
- C) Lower performance, lower periodicity, higher memory requirements
- D) Lower performance, higher periodicity, lower memory requirements
- E) Higher performance, lower periodicity, lower memory requirements

Choosing your RNG – Poll 4

- How does MT compare wrt. MINSTD?
 - Higher performance: lower time to generate 1 RN
- A) Higher performance, higher periodicity, higher memory requirements
- **B) Lower performance, higher periodicity, higher memory requirements**
- C) Lower performance, lower periodicity, higher memory requirements
- D) Lower performance, higher periodicity, lower memory requirements
- E) Higher performance, lower periodicity, lower memory requirements

Choosing your RNG – Poll 5

- For a certain random sequence, in the past we drew the number 189 immediately after the number 25. Now in the same sequence, we just drew the number 25 again.
- A) MT: the next number will be 189. LCG: the next number will be 189
- B) MT: the next number will be random. LCG: the next number will be 189
- C) MT: the next number will be 189. LCG: the next number will be random
- D) MT: the next number will be random. LCG: the next number will be random

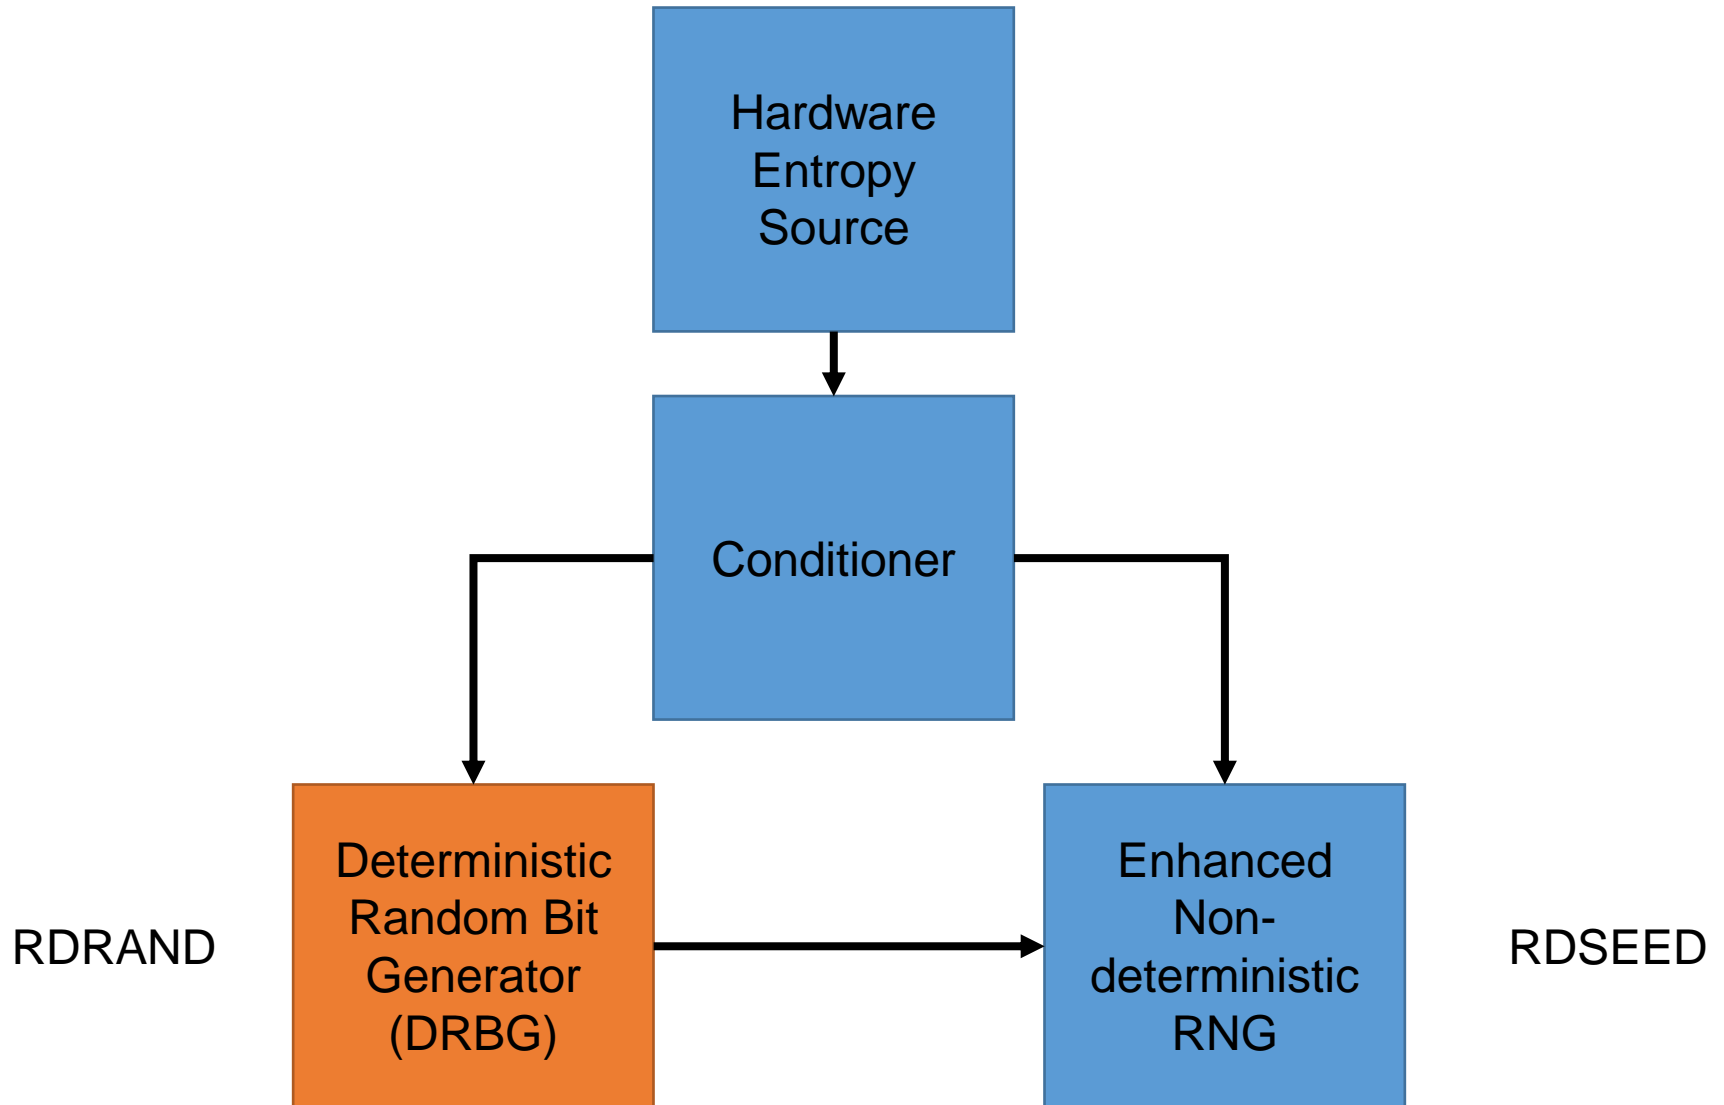
Choosing your RNG – Poll 5

- For a certain random sequence, in the past we drew the number 189 immediately after the number 25. Now in the same sequence, we just drew the number 25 again.
- A) MT: the next number will be 189. LCG: the next number will be 189
- **B) MT: the next number will be random. LCG: the next number will be 189**
- C) MT: the next number will be 189. LCG: the next number will be random
- D) MT: the next number will be random. LCG: the next number will be random

Quiz Q4 – Poll 6

- Is it possible to generate true random numbers from a digital computer (e.g. x86 architecture)?
 - Yes
 - No

Intel DRNG



Entropy Pool

- Entropy pool: where you **store the true randomness**
 - E.g.: hardware, /dev/random
- Depletion:
 - Every time you use the pool to generate a RN, you gain information about it
 - The next query might not be random!
 - The pool usually “knows” it is depleted -> needs **more randomness** (time) to replenish

RDRAND - Example

```
#include <iostream>

int main() {
    unsigned long long int random;



    for (int i = 0; i < 10; ++i) {
        if (__builtin_ia32_rdrand64_step(&random)) break;
    }

    std::cout << random << std::endl;
}
```

- GCC specific: compilation with `g++ -mrdrnd`
- Available in Intel Ivy Bridge and AMD Zen onwards
- Downside: lack of seeding means **losing reproducibility**

Generating Other Distributions

- The basic algorithms generate a uniform $u \in [0,1)$
- What if we need **other distributions**?
 - Common example: **normal (Gaussian) distribution** $N(\mu, \sigma)$
 - $N(\mu, \sigma) = \sigma * N(0,1) + \mu$
- We would like a map $u \rightarrow z \in N(0,1)$
- Very **common trick**: map $u_1, u_2 \in [0,1) \rightarrow z_1, z_2 \in N(0,1)$



Normal Distribution - Box-Muller Algorithm

- Draw 2 uniform samples u_1 and u_2

$$z_1 = \sqrt{-2\ln(u_1)} \cos(2\pi u_2)$$

$$z_2 = \sqrt{-2\ln(u_1)} \sin(2\pi u_2)$$

```
import matplotlib.pyplot as plt
import numpy as np
```


```
u1 = np.random.random(10000)
u2 = np.random.random(10000)
```

```
z1 = np.sqrt(-2*np.log(u1))*np.cos(2*np.pi*u2)
z2 = np.sqrt(-2*np.log(u1))*np.sin(2*np.pi*u2)
```

```
fig, ax = plt.subplots()
ax.hist(z1, 25, density=1, label="z1")
ax.hist(z2, 25, density=1, label="z2")
ax.legend()
```

```
plt.savefig('boxmuller.png', dpi=300)
```

Normal Distribution - Box-Muller Algorithm



Random Numbers in C++11

- C++11 introduced the `<random>` header
 - <https://www.youtube.com/watch?v=6DPkyvkMkk8>
- Random number engines and engine adaptors
 - E.g. `std::mersenne_twister_engine`
 - These folks mess around with the constants and the internal state of the RNGs
 - Don't mess around with the given constants
- Predefined RNGs
 - E.g. `std::mt19937_64` is an instantiation of the engine with the correct constants for a 64-bit random number
 - A collection of popular RNGs including both 32 and 64-bit Mersenne Twister

Random Numbers in C++11

- Random Number Distributions
 - E.g. `std::uniform_int_distribution`
 - It can be imagined as a view of the RNG
 - Given a certain generator, **extract a new random number according to the distribution**
 - Many common choices, **including Gaussian, Poisson...**
- Non-deterministic Random Numbers
 - `std::random_device` generates a non-deterministic uniform 32-bit int
 - It **may** use hardware-specific RNG.
 - Does your implementation actually use RDRAND? `^_(ツ)_/`
 - Common application: **generate a good seed for your RNG**
 - It can be used for general purpose, but might become slower

Random Numbers in C++11 - Example

```
#include <iostream>
#include <random>

int main()
{
    std::random_device rd;
    unsigned int rd_seed = rd();
    std::mt19937_64 gen(rd_seed);
    std::normal_distribution<double> d(0,1);

    for(int i=0; i<5; ++i) {
        std::cout << d(gen) << std::endl;
    }
}
```

Poll 7

- In which line might there be a problem?

```
#include <iostream>
#include <random>
```

```
int main()
```

```
{
```

```
    std::random_device rd;
```

```
    unsigned int rd_seed = rd();
```

```
    std::mt19937_64 gen(rd_seed);
```

```
    std::normal_distribution<double> d(0,1);
```

```
    for(int i=0; i<5; ++i) {
```

```
        std::cout << d(gen) << std::endl;
```

```
    }
```

```
}
```

• A

• B

• C

• D

• E

Poll 7

- In which line might there be a problem?

```
#include <iostream>
#include <random>
```

```
int main()
```

```
{
```

```
    std::random_device rd;
```

```
    unsigned int rd_seed = rd();
```

```
    std::mt19937_64 gen(rd_seed);
```

```
    std::normal_distribution<double> d(0,1);
```

```
    for(int i=0; i<5; ++i) {
```

```
        std::cout << d(gen) << std::endl;
```

```
    }
```

```
}
```

• A

• B

• C

• D

• E

Random Numbers in a Parallel System – Poll 7

- What happens when different RNGs are initialized in a parallel system with the same seed?
- A) Each stream generates different, independent RNs
- B) The result of each stream is random, depending on initialization
- C) All threads read from the same RNG, leading to data races
- D) All streams generate the same sequence of RNs

Random Numbers in a Parallel System – Poll 7

- What happens when different RNGs are initialized in a parallel system with the same seed?
- A) Each stream generates different, independent RNs
- B) The result of each stream is random, depending on initialization
- C) All threads read from the same RNG, leading to data races
- **D) All streams generate the same sequence of RNs**

Random Numbers in a Parallel System

- Random number server
 - The program only has **a single RNG** and all threads and processes draw numbers from it
 - Commonly used for GPUs: the CPU generates a list of random numbers which is then loaded to the GPU
- Per-stream generator
 - Assign a **different RNG per thread** or process
 - **Pay attention to the seeds!**
 - One approach: 1 master generator with 1 seed -> 1 seed per process/thread

Sidenote: CS RNG

- Working definition [Kneusel] of Cryptographically Safe RNG:
 - A CSPRNG is a random number generator that passes the next-bit test and one where an attacker's knowledge of the state of the generator at time t makes knowing the state at any previous time impossible.
- Irrelevant for simulations
- But a requirement for cryptography!
- MT is NOT cryptographically safe

Rules of Thumb

- If you do not want to think, use the **Mersenne Twister**
 - Consider the 64-bit version when generating doubles...
- Each thread or process needs to have **its own RNG**
 - With **different seed!**
- Avoid LCG if possible
 - Every time you use C-style `rand()`, a transistor cries
 - In general, their period is too short for large scale simulations
- Hardware RNGs might make sense in many applications!
 - Watch out for **entropy pool depletion**
- Correct seeding can be hard – but allows for reproducibility!

Monte Carlo


- What is/Where is Monte Carlo?
- Term dates to the **Manhattan Project**
- Very loosely defined: any type of **algorithm which uses random sampling to achieve a numerical result**
 - The problem itself might be deterministic!
 - Las Vegas methods (subclass): uses randomness but achieves a deterministic result
- “Monte Carlo” has in practice many definitions for each field

Quiz – Question 5

- How would you estimate the value of π using random numbers?

Monte Carlo Integration – one way

- Draw a circle inside the unit square
 - $g(x) = \sqrt{1 - x^2}$
- Generate pairs of random numbers
 - $x, y \in [0, 1)$
- Count the number of hits inside
 - $y < g(x)$
- Calculate the ratio:
 - $\frac{\pi}{4} \approx \frac{\#_{inside}}{\#_{total}}$



Monte Carlo Integration – another way

$$I = \int_D g(\mathbf{x}) d\mathbf{x}$$



$$\hat{I}_m = \frac{1}{m} \{g(\mathbf{x}^{(1)}) + \dots + g(\mathbf{x}^{(m)})\}$$

- Where $\mathbf{x}^{(n)}$ are drawn uniformly from D
- For the circle, $D: x \in [0,1)$ and $g(x) = \sqrt{1 - x^2}$

Take-home message – version 2

- Be **careful with your choice** of RNG
 - The Mersenne Twister is the standard for a reason
 - Hardware RNGs are becoming more useful
 - Older LCG are to be avoided whenever possible
- **Seeding** is a consequence of pseudo-RNG
 - Not always a bad thing – reproducibility!
- Watch out for pitfalls on parallel systems
 - Usually, you must make sure that each process/thread is generating a **different random number sequence**

Quiz

- Q1: Consider a 32-bit LCG. How many random numbers can you generate before the sequence repeats?
- Q2: What are the consequences of entropy pool depletion? How would you handle it?
- Q3: What are the requirements to initialize correctly one RNG per thread?
- Q4: What is a different between a process and a thread? Which resources can be shared among them?
- Q5: What does the directive `#pragma omp atomic` do?

Next stop

- Shared Memory Parallel Optimization
- OpenMP