

Computational Science on Many-Core Architectures

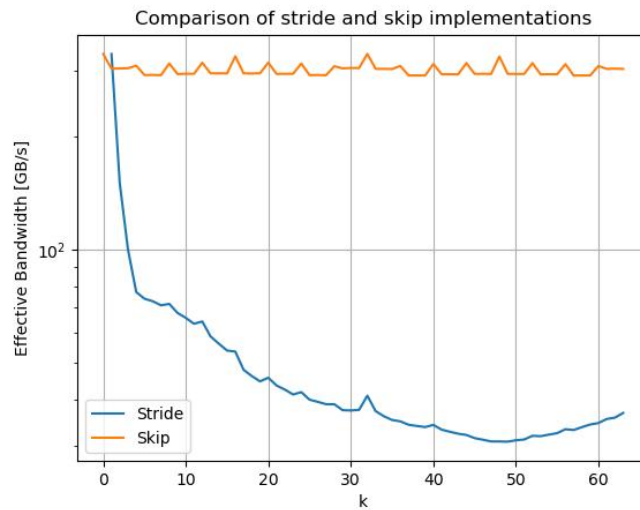
Exercise 3

Leon Schwarzäugl

3. November 2023

The code for all tasks can be found at: https://github.com/Swarsel/CSE_TUWIEN/tree/main/WS2023/Many-Core%20Architectures/e3

1 Strided and Offset Memory Access



1.1 a)

As we can see, the effective bandwidth drops very quickly with rising k . This can be explained due to data locality: data is accessed in a contiguous manner by hardware, which we are not making use of when performing large strides, as

then our threads are accessing addresses that are far apart. Hence the effective bandwidth is higher for small k . It should be recommended to keep this effect in mind and implement code in a manner that makes use of data locality - in this case, avoiding larger strides.

1.2 b)

Here we see a more constant effective bandwidth; this is again to be expected, since we are simply leaving a contiguous block of addresses out at the start; the fluctuations can be explained by the same effect, as the effective bandwidth seems to be especially high for multiples of 4. It seems that leaving away a number of entries that is not divisible by it causes the need to load another chunk of data which in turn lessens the effective bandwidth a little. As for a recommendation, it can be advised to keep this effect in mind and only skip blocks of 4 entries at once, for example.

2 Dense Matrix Transpose

2.1 a)

I identified two issues:

1. There is a memory leak, since the `free` and `cudaFree` functions are not called at the end.
2. The kernel only works up to the entry assigned to the n -th thread, as there is nothing implemented towards looping over the rest of the data.

2.2 b)

A simple kernel was implemented that stores the upper triangular value in a temporary variable and then updates both values corresponding to the row-column pair:

```
__global__ void transpose(double *A, int N)
{
    int thread_id = blockIdx.x*blockDim.x + threadIdx.x;
    unsigned int total_threads = blockDim.x * gridDim.x;

    for (unsigned int i = thread_id; i < N*N; i += total_threads) {
        int r = i / N;
        int c = i % N;

        if (r < c) {
            double temp = A[N * r + c];
            A[N * r + c] = A[N * c + r];
            A[N * c + r] = temp;
        }
    }
}
```

Also, the mentioned free-operations were added.

For the given $N = 10$ an effective bandwidth of $0.106667 \frac{\text{GB}}{\text{sec}}$ was measured.

2.3 c)

Dropping the in-place requirement simplifies the value-updating and we can drop the need for the temporary variable while now again updating for every value a single entry:

```
__global__ void transpose(double *A, double *B, int N)
{
    int thread_id = blockIdx.x*blockDim.x + threadIdx.x;
    unsigned int total_threads = blockDim.x * gridDim.x;

    for (unsigned int i = thread_id; i < N*N; i += total_threads) {
        int r = i / N;
        int c = i % N;

        if (r < N && c < N) {
            B[N * r + c] = A[N * c + r];
        }
    }
}
```

2.4 d)

The implemented kernel is close to verbatim to the one presented in the NVIDIA-blog, while only changing to the tile sizes wanted in our case:

```
__global__ void transpose(double *A)
{
    __shared__ float tile[16][16];

    int x = blockIdx.x * 16 + threadIdx.x;
    int y = blockIdx.y * 16 + threadIdx.y;
    int width = gridDim.x * 16;

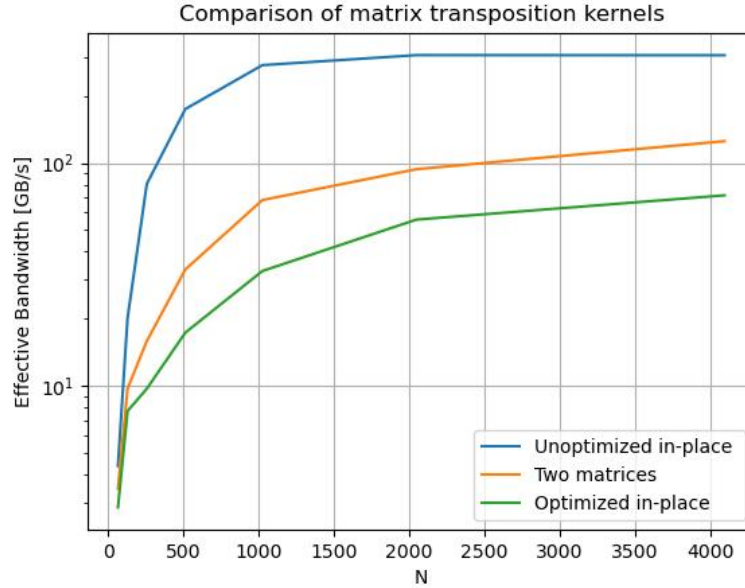
    for (int j = 0; j < 16; j += 8)
        tile[threadIdx.y+j][threadIdx.x] = A[(y+j)*width + x];

    __syncthreads();

    x = blockIdx.y * 16 + threadIdx.x;
    y = blockIdx.x * 16 + threadIdx.y;

    for (int j = 0; j < 16; j += 8)
        A[(y+j)*width + x] = tile[threadIdx.x][threadIdx.y + j];
}
```

2.5 e)



The comparison of the different kernels shows that our optimizations have had a positive impact on the effective bandwidth. We can see that the dropping of the in-place requirement led to a first increase in performance. The implementation using shared memory where the reads and writes are coalesced leads to the biggest increase in effective bandwidth.

3 Trick AND Treat

At the start, I just moved the lines `timer.reset()` and `timer.get()` to convenient spots, which already showed quite well that the GPU version was by far superior.

Then I thought I should make extra sure and added a little sanity test to the CPU version, because it performed so poorly when compared to the GPU version. Even for low $N \approx 10000$, the CPU version already took close to the maximum allotted 30 seconds.

To make doubly sure, I ran some more tests on the same N for the GPU version, where surprisingly I noticed that the runtime t is about $t = -0.000002s$ - so it literally saves time to run that code. :-)