

Computational Science on Many-Core Architectures

Exercise 7

Leon Schwarzäugl

December 5, 2023

The code for all tasks can be found at: https://github.com/Swarsel/CSE_TUWIEN/tree/main/WS2023/Many-Core%20Architectures/e7

1 Pipelined Conjugate Gradients

I implemented the pipelined conjugate gradients using the two kernels introduced in the lecture:

```
__global__ void XRP(double *solution, double *p, double* r, double* Ap, double* result,
                   int N, double alpha, double beta) {
    double dot{0};
    for (int i = blockIdx.x * blockDim.x + threadIdx.x; i < N; i += blockDim.x * gridDim.x){
        solution[i] += alpha * p[i];
        r[i] -= alpha * Ap[i];
        p[i] = r[i] + beta * p[i];
        dot += r[i] * r[i];
    }
    for (int j=warpSize/2; j>0; j=j/2) {
        dot += __shfl_xor_sync(0xffffffff, dot, j);
    }

    if (threadIdx.x % warpSize == 0) {
        atomicAdd(result, dot);
    }
}

__global__ void API(int N, int *csr_rowoffsets, int *csr_colindices, double *csr_values,
                   double *p, double *Ap, double *result_pAp, double *result_ApAp) {
    double pAp{0}, ApAp{0};
    for (int row = blockIdx.x * blockDim.x + threadIdx.x; row < N; row += blockDim.x * gridDim.x) {
        double sum = 0;
        for (int jj = csr_rowoffsets[row]; jj < csr_rowoffsets[row + 1]; ++jj) {
            sum += csr_values[jj] * p[csr_colindices[jj]];
        }
        Ap[row] = sum;
        pAp += p[row] * Ap[row];
        ApAp += Ap[row] * Ap[row];
    }

    for (int j=warpSize/2; j>0; j=j/2) {
        pAp += __shfl_xor_sync(0xffffffff, pAp, j);
        ApAp += __shfl_xor_sync(0xffffffff, ApAp, j);
    }

    if (threadIdx.x % warpSize == 0) {
        atomicAdd(result_pAp, pAp);
        atomicAdd(result_ApAp, ApAp);
    }
}
```

The CG algorithm itself was adjusted to perform one initial calculation of α_0, β_0, Ap_0 using the ‘old’ kernels and then use the more efficient kernels for the rest of the iterations:

[...]

```

    cuda_csr_matvec_product<<<512, 512>>>(N, csr_rowoffsets, csr_colindices, csr_values, cuda_p

    cudaMemcpy(cuda_scalar, &zero, sizeof(double), cudaMemcpyHostToDevice);
    cudaMemcpy(cuda_scalar2, &zero, sizeof(double), cudaMemcpyHostToDevice);
    cuda_dot_product<<<512, 512>>>(N, cuda_p, cuda_Ap, cuda_scalar);
    cuda_dot_product<<<512, 512>>>(N, cuda_Ap, cuda_Ap, cuda_scalar2);
    cudaMemcpy(&alpha, cuda_scalar, sizeof(double), cudaMemcpyDeviceToHost);
    cudaMemcpy(&beta, cuda_scalar2, sizeof(double), cudaMemcpyDeviceToHost);
    alpha = residual_norm_squared / alpha;
    beta = alpha*alpha*beta / residual_norm_squared - 1;

    int iters = 0;
    cudaDeviceSynchronize();
    timer.reset();
    while (1) {


        cudaMemcpy(cuda_scalar, &zero, sizeof(double), cudaMemcpyHostToDevice);
        cudaMemcpy(cuda_scalar2, &zero, sizeof(double), cudaMemcpyHostToDevice);
        XRP<<<512, 512>>>(cuda_solution, cuda_p, cuda_r, cuda_Ap, cuda_scalar, N, alpha, beta);
        cudaMemcpy(&residual_norm_squared, cuda_scalar, sizeof(double), cudaMemcpyDeviceToHost);
        cudaMemcpy(cuda_scalar, &zero, sizeof(double), cudaMemcpyHostToDevice);
        API<<<512, 512>>>(N, csr_rowoffsets, csr_colindices, csr_values, cuda_p, cuda_Ap, cuda_sc
        cudaMemcpy(&alpha, cuda_scalar, sizeof(double), cudaMemcpyDeviceToHost);
        cudaMemcpy(&beta, cuda_scalar2, sizeof(double), cudaMemcpyDeviceToHost);
        alpha = residual_norm_squared / alpha;
        beta = (alpha*alpha*beta / residual_norm_squared) - 1;

```

[...]


Using what we learned this and last week, I had four different versions of (Assembly, CG) at my disposal for benchmarking. Below are the results. As always, each system size was ran 14 times, with the lowest and highest 2 times omitted before averaging over the rest.

First, the results with all four systems compared against each other:



I have to say that I am disappointed with the GPU matrix assembly results - after last weeks results, I thought I would reach way better runtimes here. But it seems not to be the case here. Interestingly, the GPU assembled implementation really takes off in runtime (in a bad way) for high N , even taking longer than the classical implementation then. Maybe I introduced a bug while transporting the pipelined implementation into the GPU assembly framework, but at least I could not find it.

Anyways, what we actually wanted to show here holds - that is, the pipelined implementations are a lot faster than the classical implementations. To get a better picture of this distinction, I ran another benchmark with more system sizes only comparing the CPU assembly versions against each other:



Depending on the size of N , we see performance increases between 10 – 40%, which is quite respectable. The best gain in that respect is received for ‘middling’ N that are neither too high nor too low.